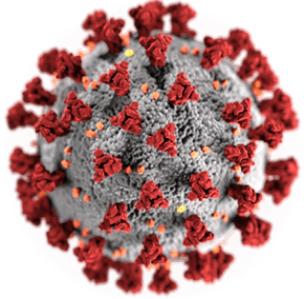


EECS 498-007 / 598-005

Deep Learning for Computer Vision

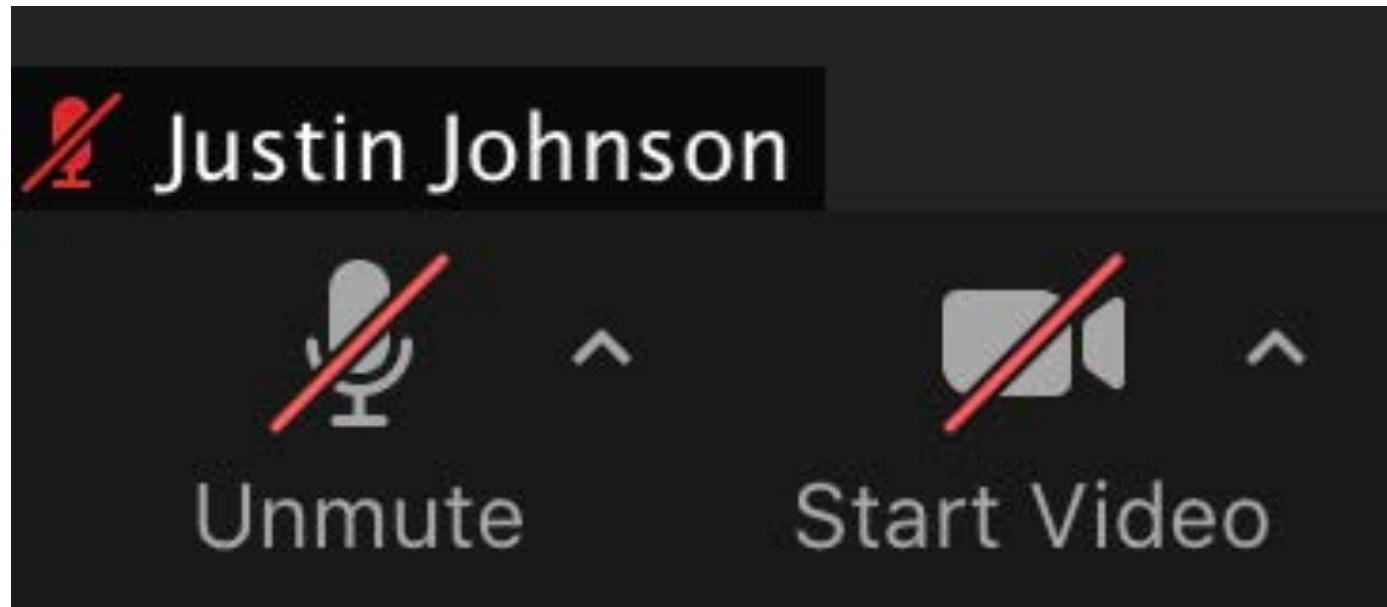
Lecture 1: Introduction

COVID-19 Edition

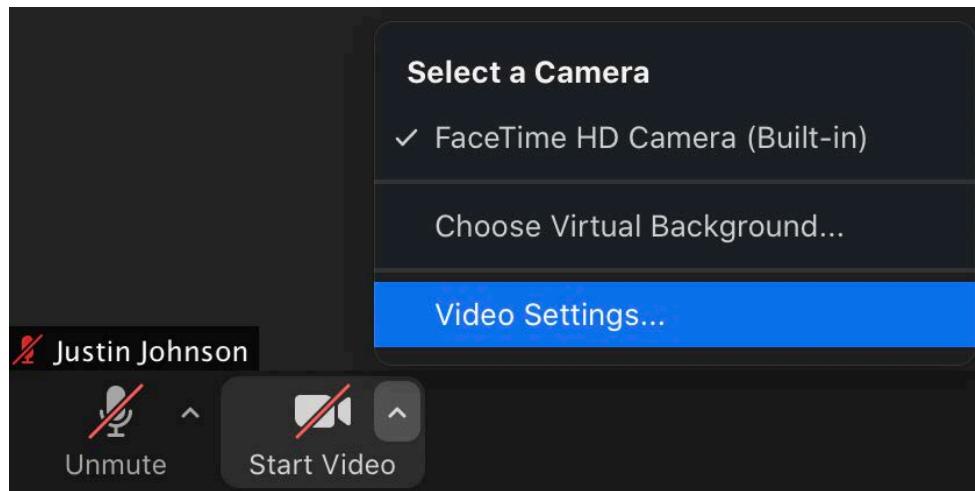


EECS 498-007 / 598-005 Deep Learning for Computer Vision Lecture 1: Introduction

Zoom Logistics: Muted, Camera Off



Zoom Logistics: Hide non-video participants



Camera: FaceTime HD Camera (Built-in)

16:9 (Widescreen) Original ratio

My Video:

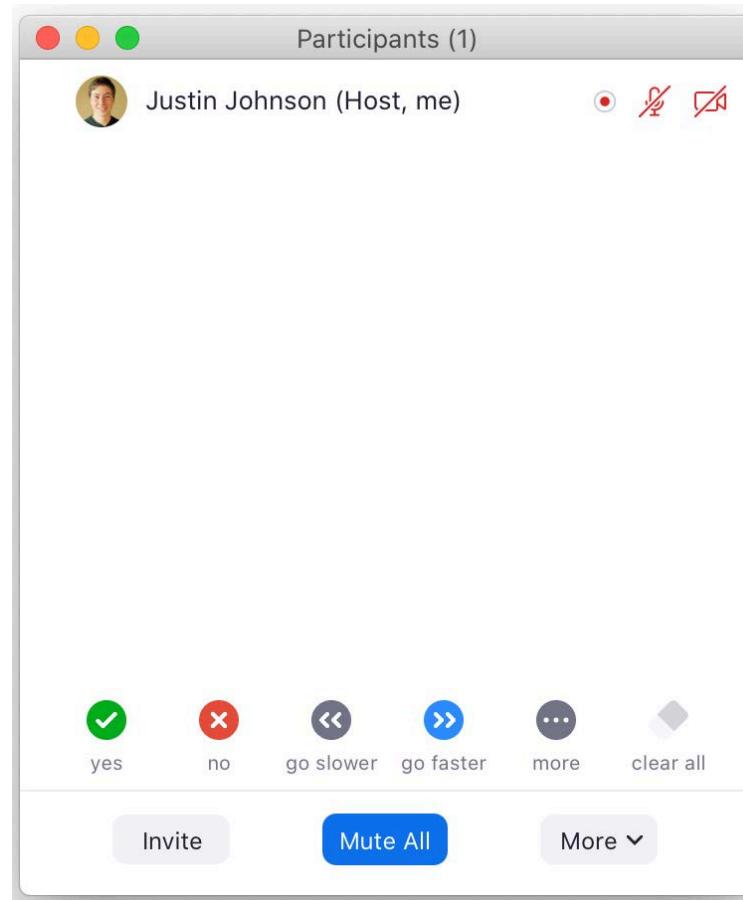
- Enable HD
- Mirror my video
- Touch up my appearance

Meetings:

- Always display participant name on their videos
- Turn off my video when joining a meeting
- Always show video preview dialog when joining a video meeting
- Hide non-video participants
- Spotlight my video when speaking
- Display up to 49 participants per screen in Gallery View

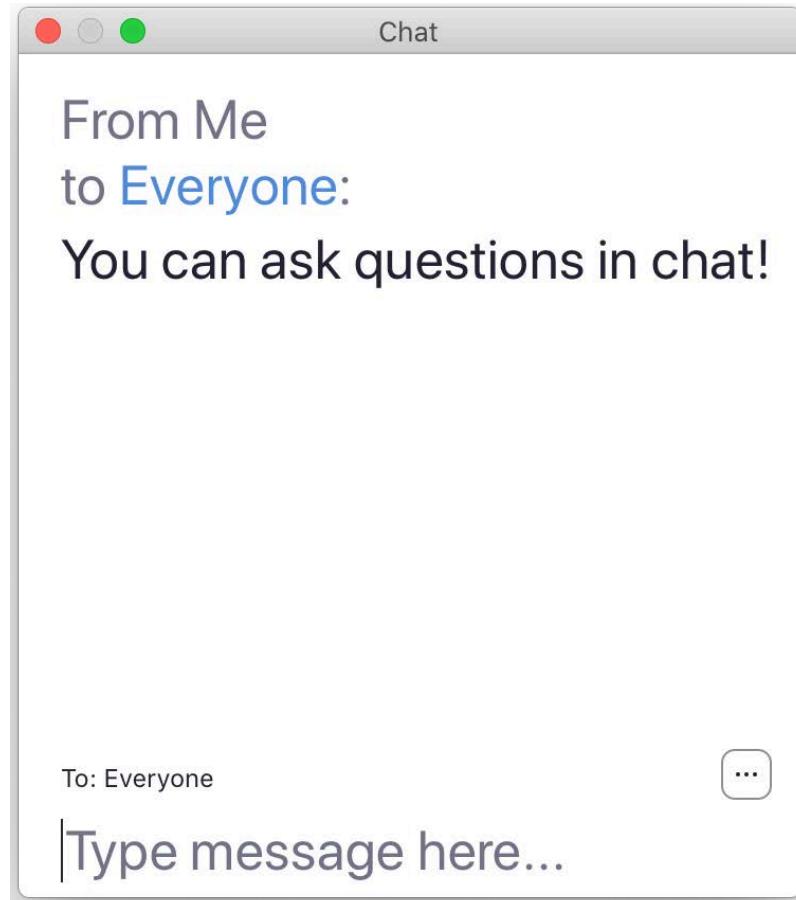
Zoom Logistics: Quick Poll

- You can use responses in “participants” panel for quick polls

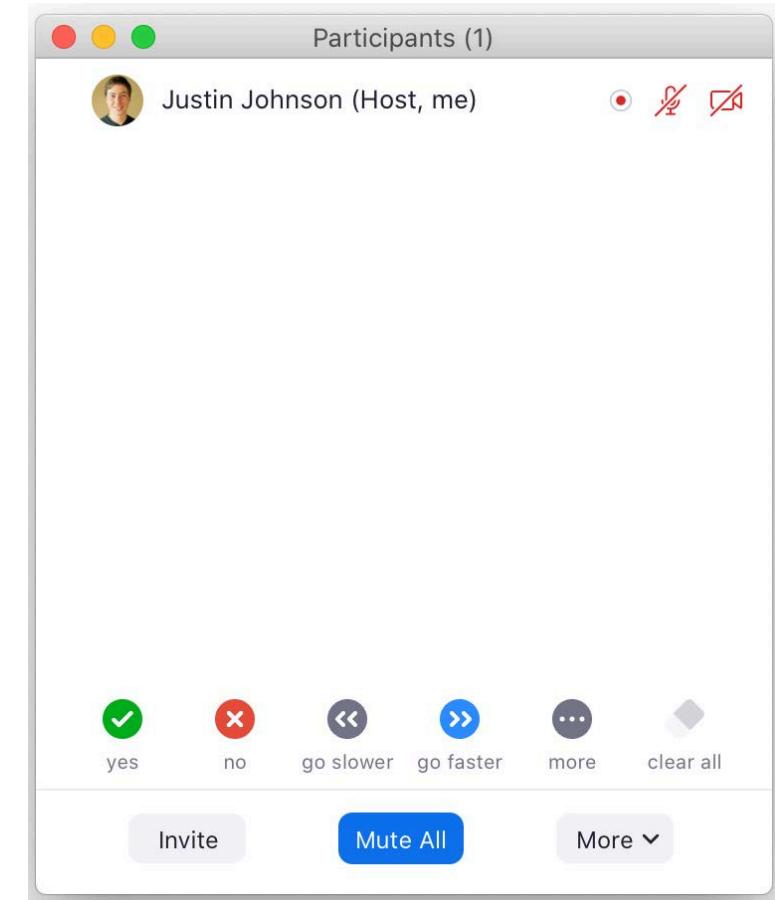


Zoom Logistics: Asking Questions

Option 1: Ask questions in chat



Option 2: "Raise Hand" in participants



Deep Learning for Computer Vision

Deep Learning for Computer Vision

Building artificial systems
that process, perceive, and
reason about visual data

Computer Vision is everywhere!



Left to right:
[Image by Roger H Goun](#) is licensed under CC BY 2.0
[Image](#) is CCO 1.0 public domain
[Image](#) is CCO 1.0 public domain
[Image](#) is CCO 1.0 public domain



Left to right:
[Image is free to use](#)
[Image](#) is CCO 1.0 public domain
[Image by NASA](#) is licensed under CC BY 2.0
[Image](#) is CCO 1.0 public domain



Bottom row, left to right:
[Image](#) is CCO 1.0 public domain
[Image by Derek Keats](#) is licensed under CC BY 2.0; changes made
[Image](#) is public domain
[Image](#) is licensed under CC-BY 2.0; changes made

Deep Learning for Computer Vision

Building artificial systems that
learn from data and experience

Deep Learning for Computer Vision

Hierarchical learning algorithms
with many “layers”, (very) loosely
inspired by the brain

Artificial Intelligence

Artificial Intelligence

Machine Learning

Computer
Vision

Artificial Intelligence

Computer
Vision

Machine Learning

Deep
Learning

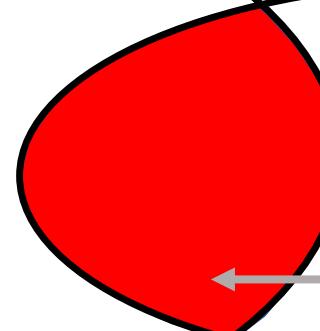
Artificial Intelligence

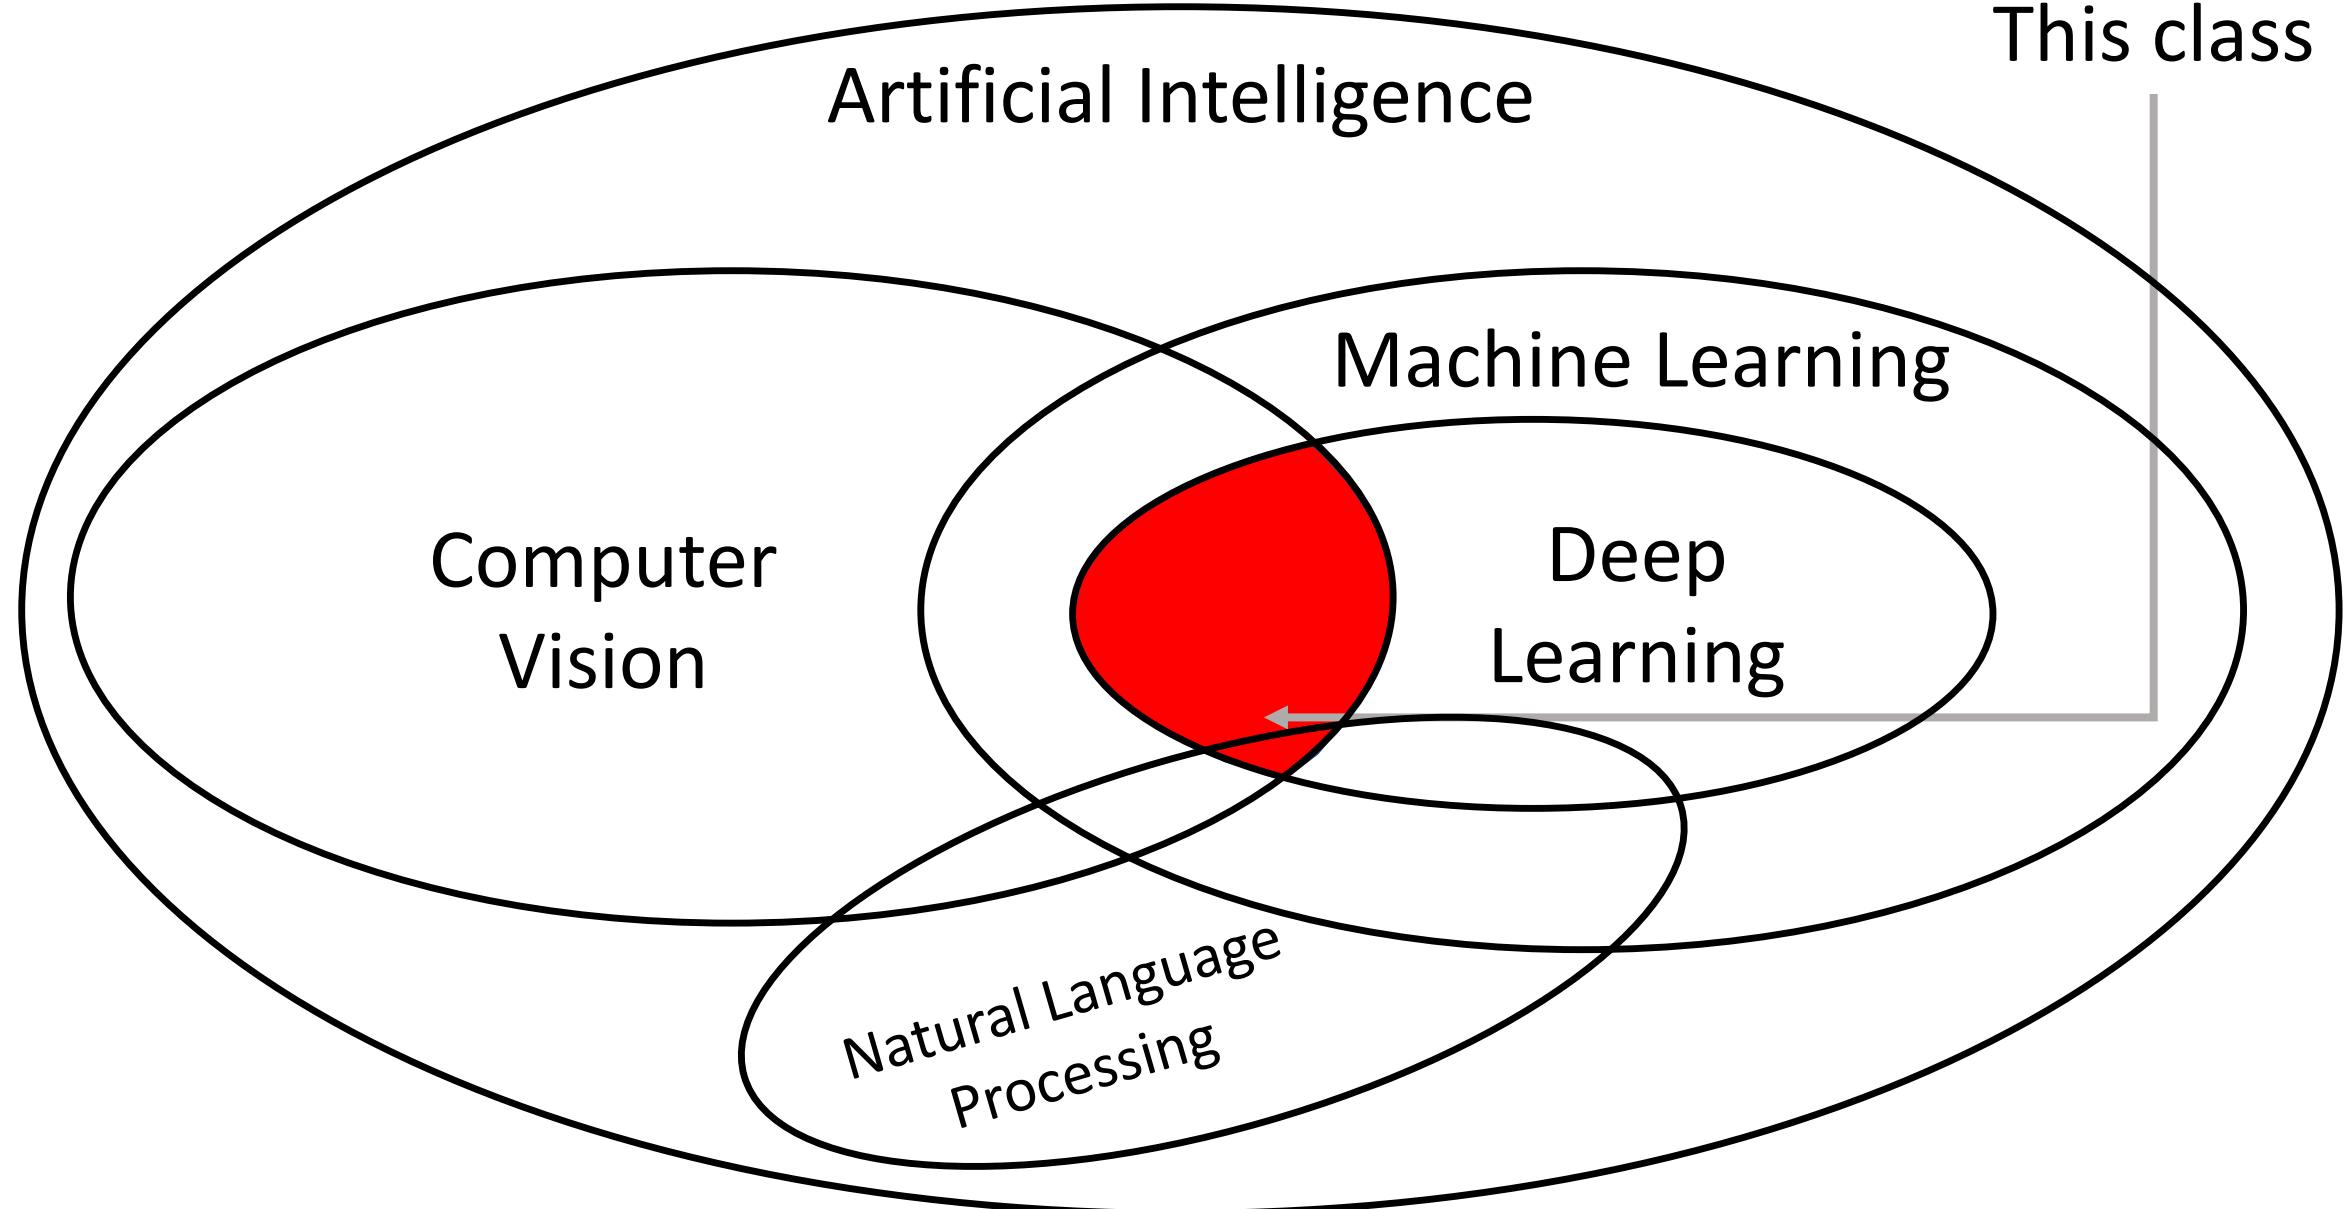
This class

Computer
Vision

Machine Learning

Deep
Learning





Artificial Intelligence

This class

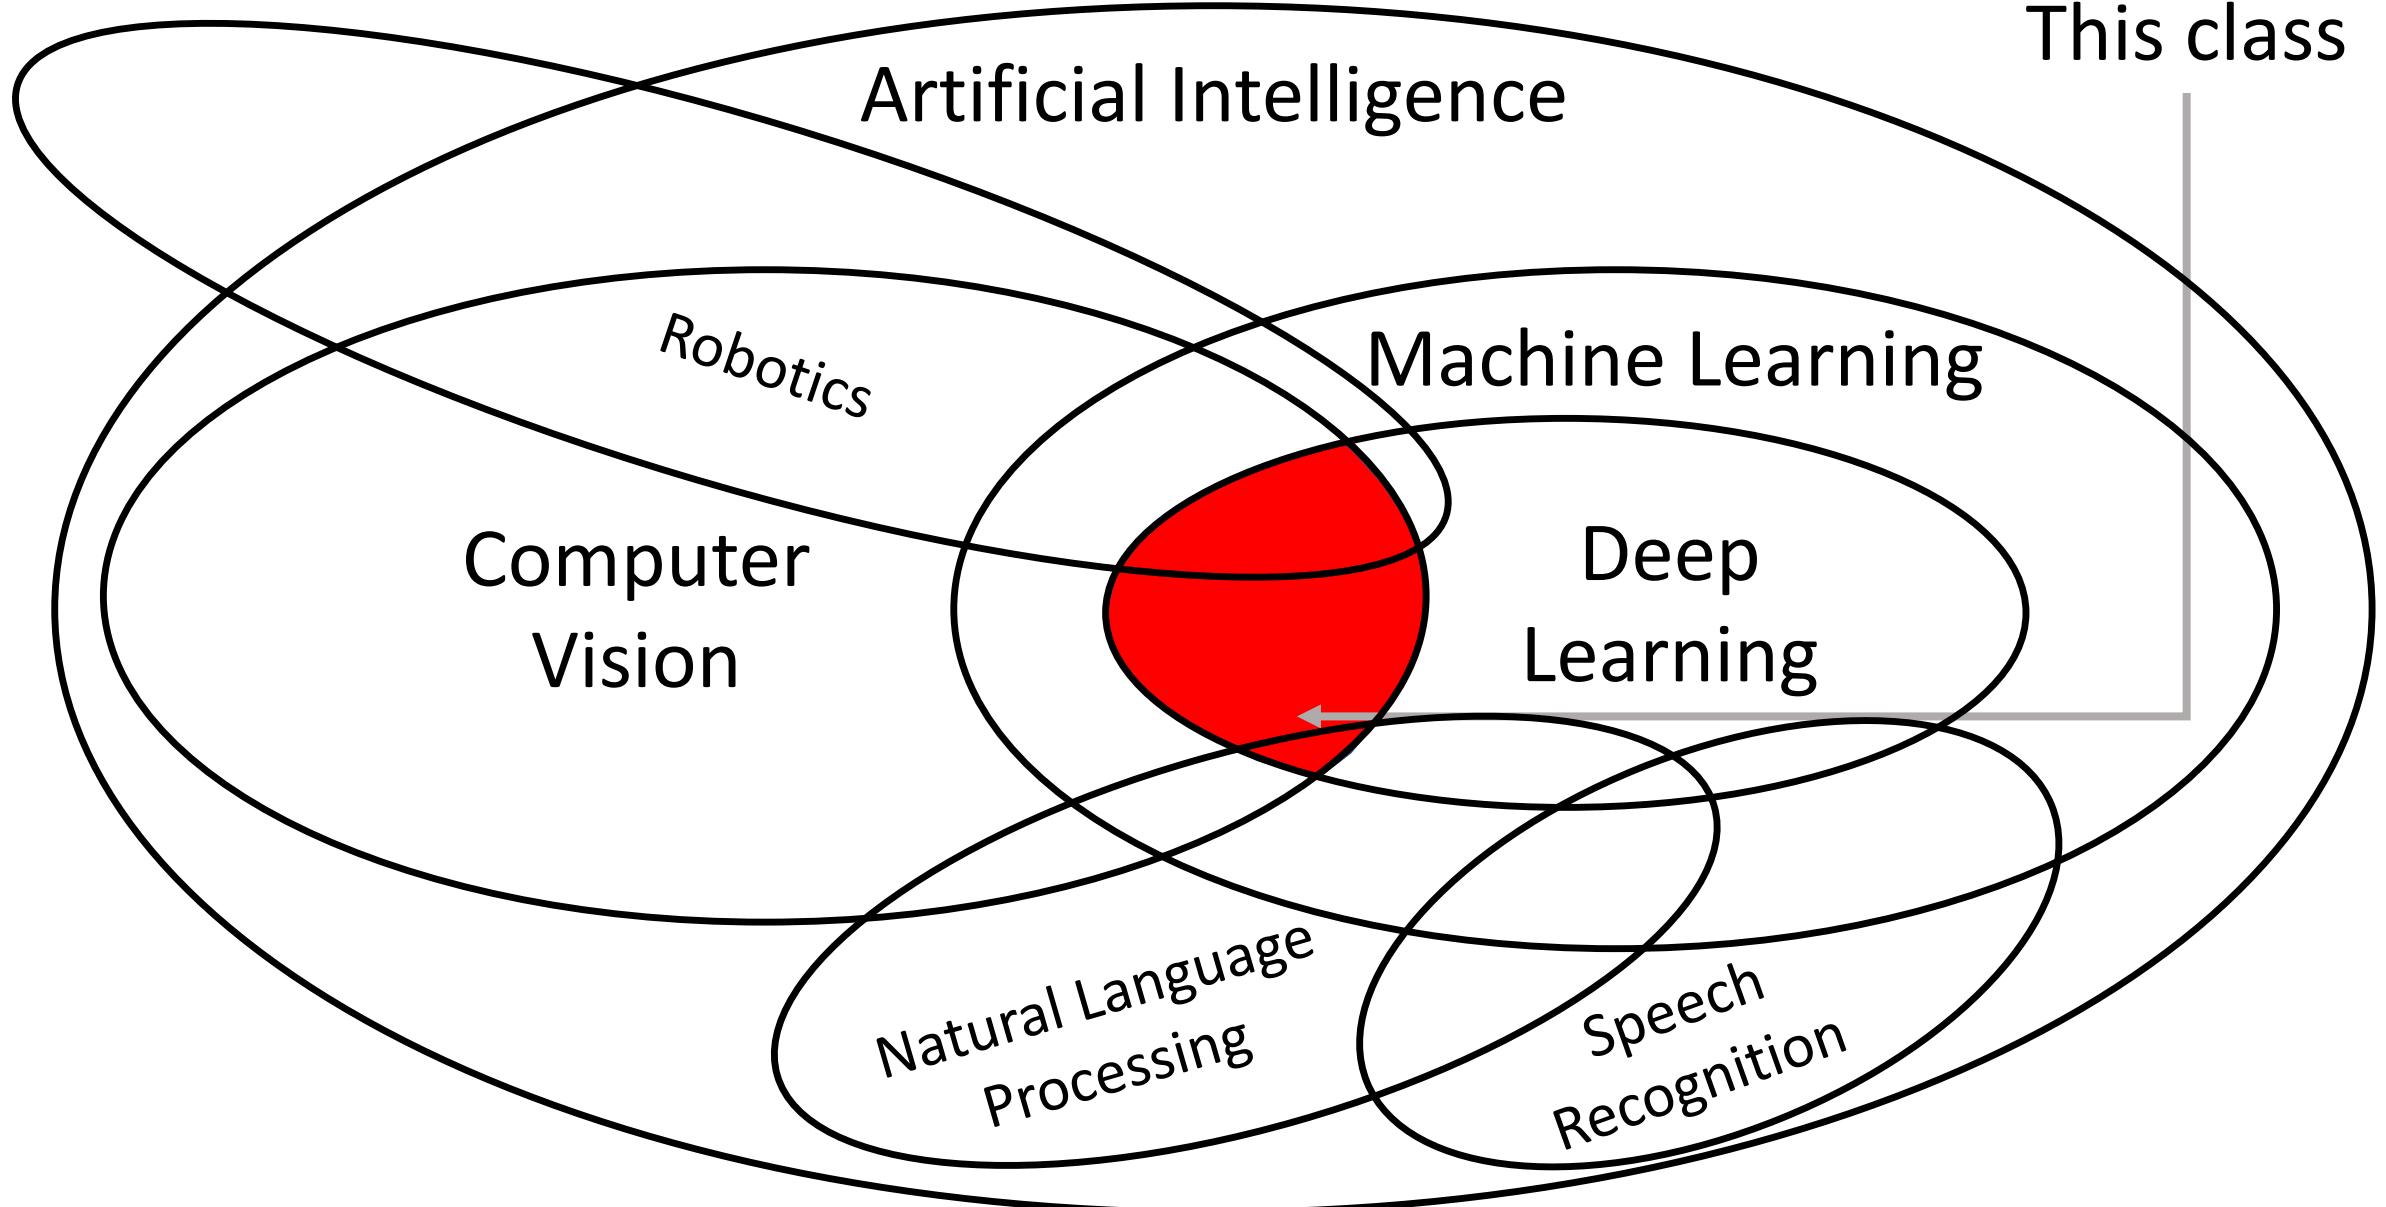
Computer
Vision

Machine Learning

Deep
Learning

Natural Language
Processing

Speech
Recognition

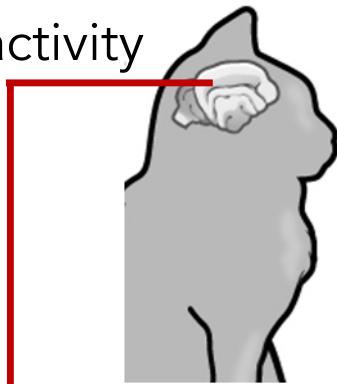


Today's Agenda

- A brief history of computer vision and deep learning
- Course overview and logistics

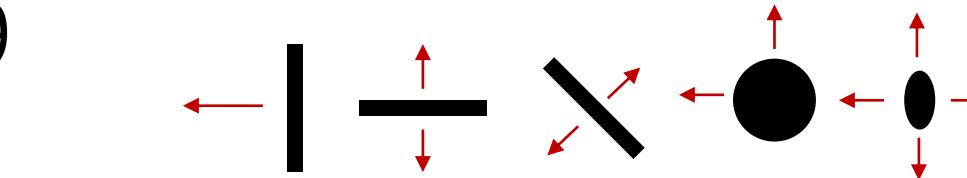
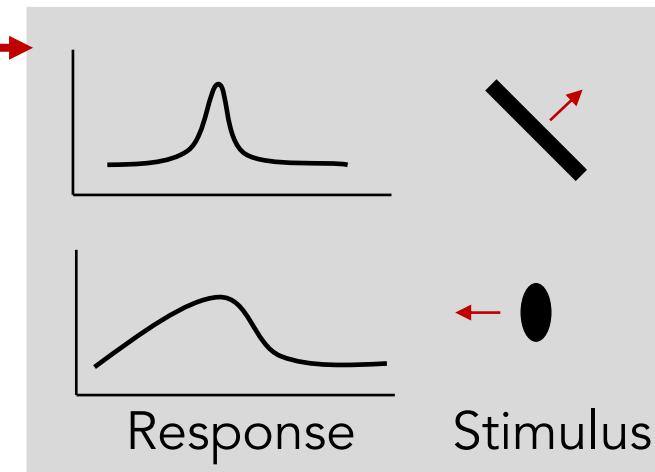
Hubel and Wiesel, 1959

Measure
brain activity

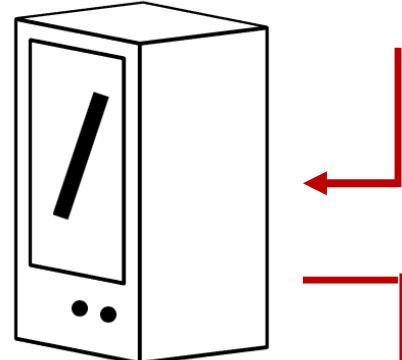


Cat image by CNX OpenStax is licensed under CC BY 4.0; changes made

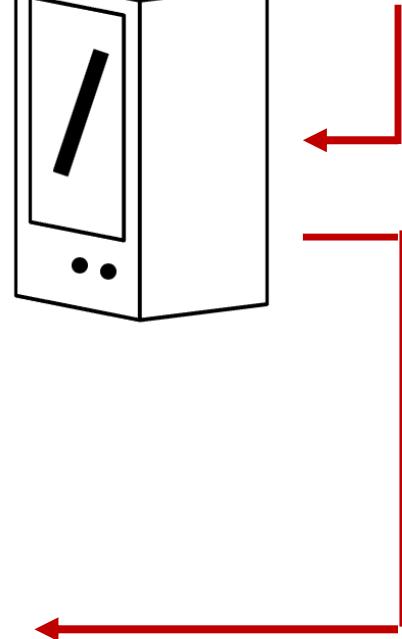
1959
Hubel & Wiesel



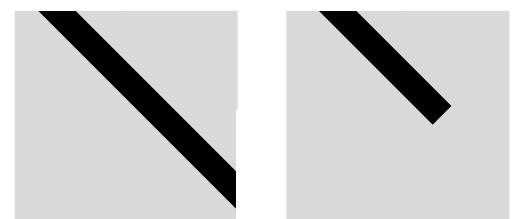
Simple cells:
Response to light orientation



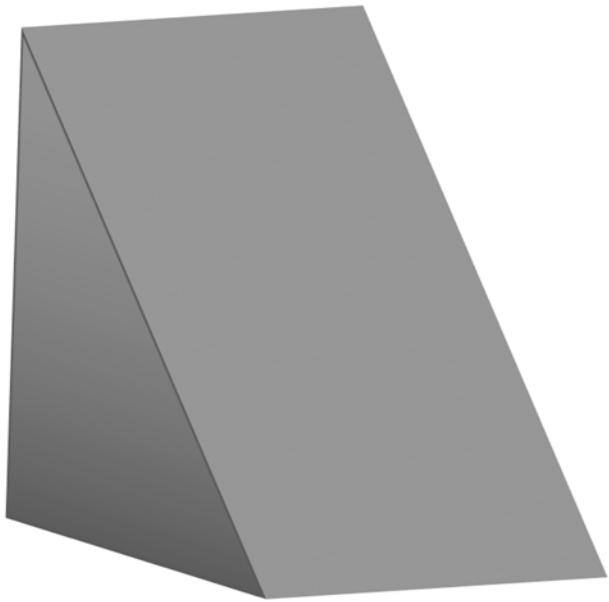
Complex cells:
Response to light orientation and movement



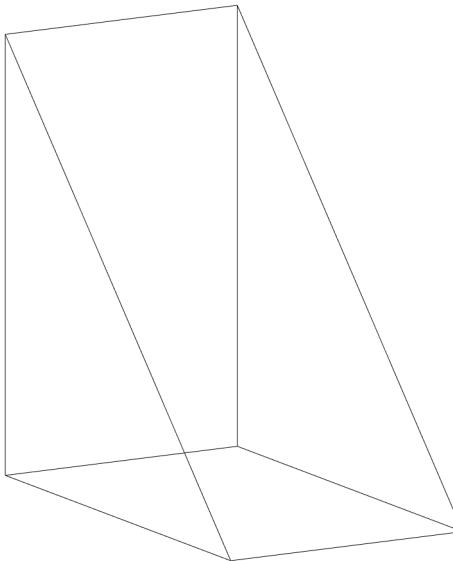
Hypercomplex cells:
response to movement with an end point



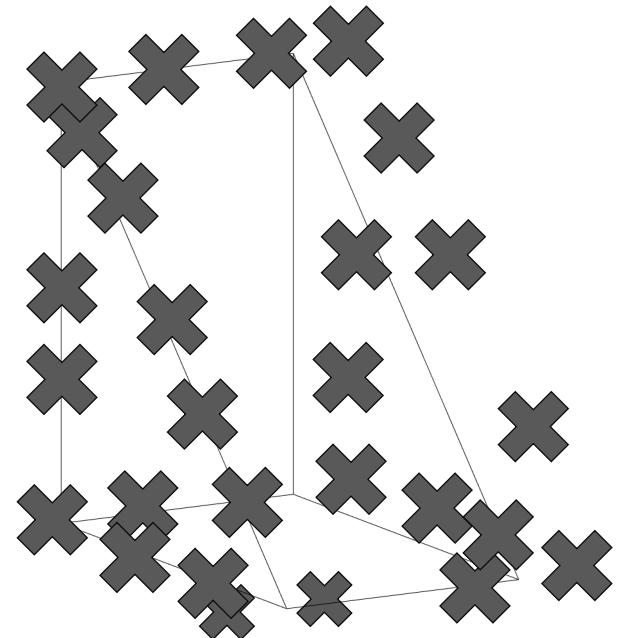
Larry Roberts, 1963



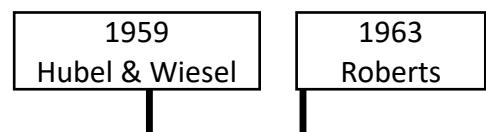
(a) Original picture



(b) Differentiated picture



(c) Feature points selected



MASSACHUSETTS INSTITUTE OF TECHNOLOGY

PROJECT MAC

Artificial Intelligence Group
Vision Memo. No. 100.

July 7, 1966

THE SUMMER VISION PROJECT

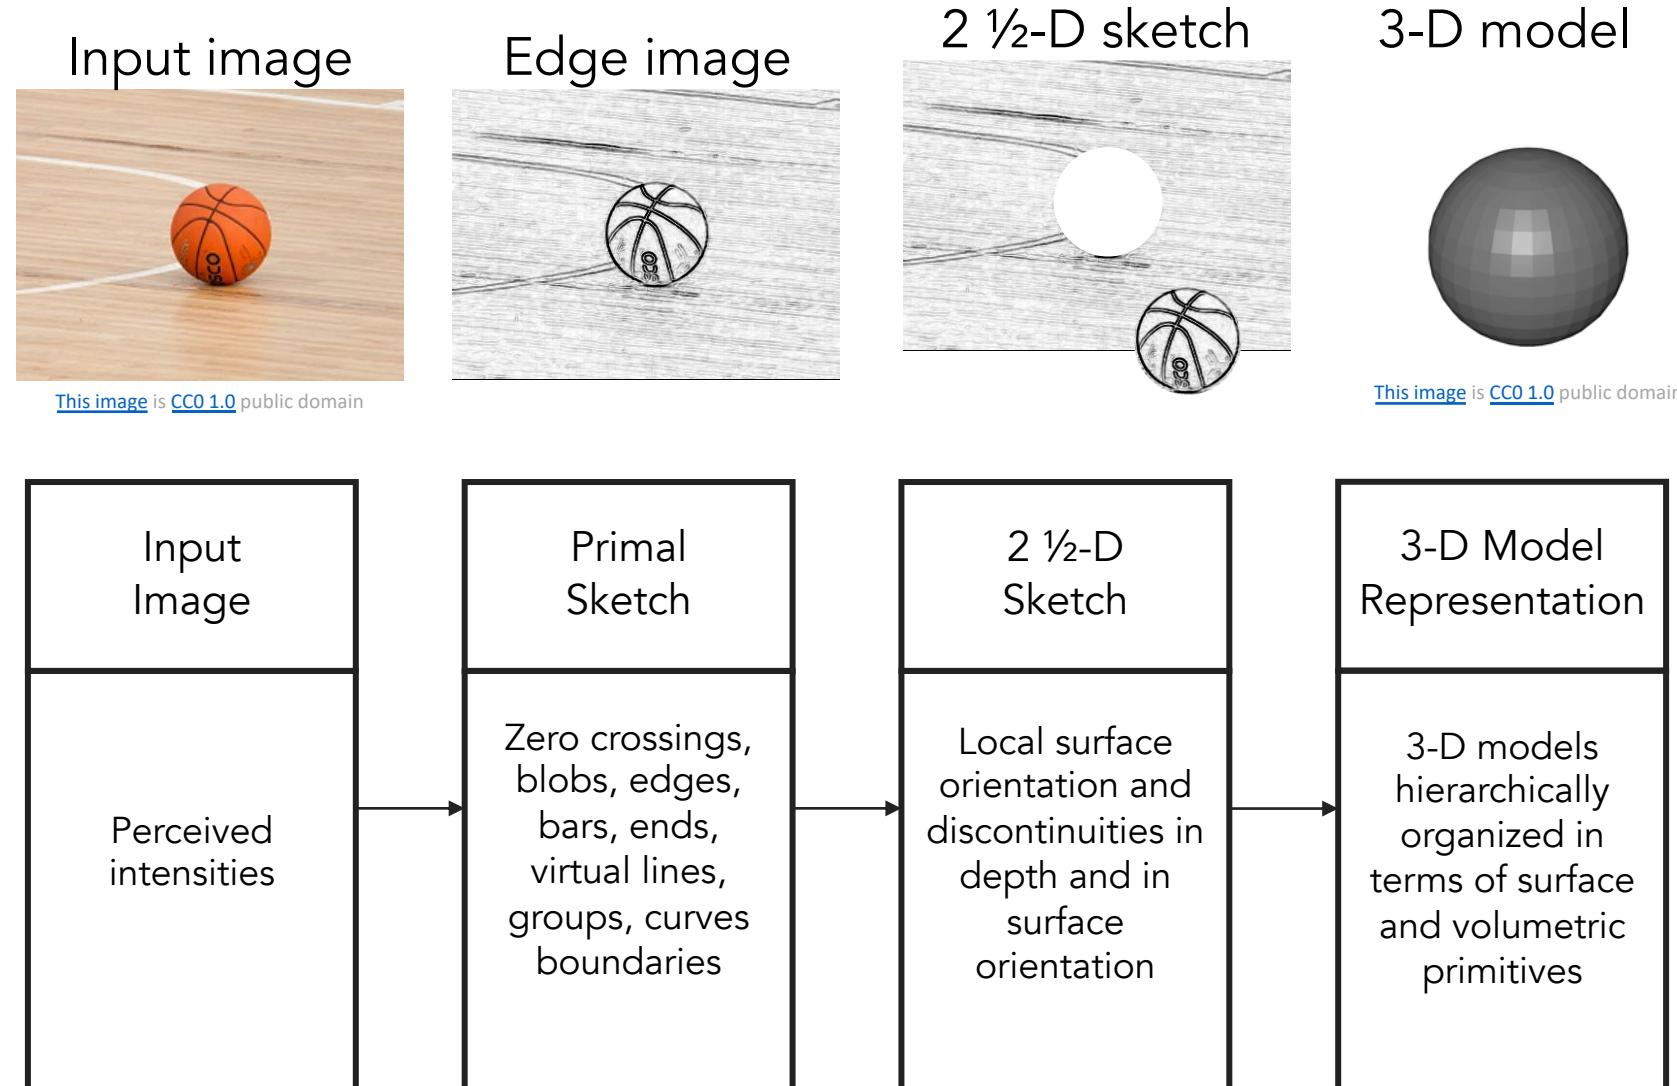
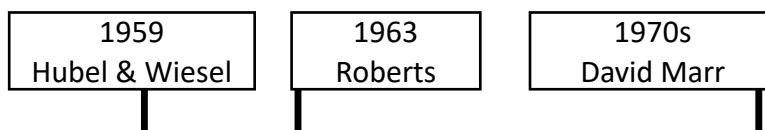
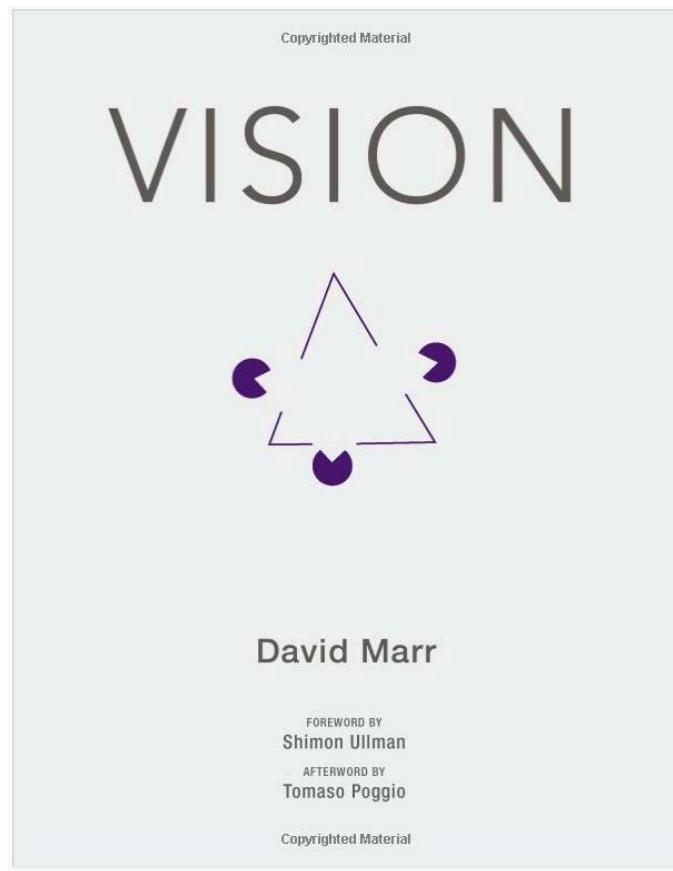
Seymour Papert

The summer vision project is an attempt to use our summer workers effectively in the construction of a significant part of a visual system. The particular task was chosen partly because it can be segmented into sub-problems which will allow individuals to work independently and yet participate in the construction of a system complex enough to be a real landmark in the development of "pattern recognition".

1959
Hubel & Wiesel

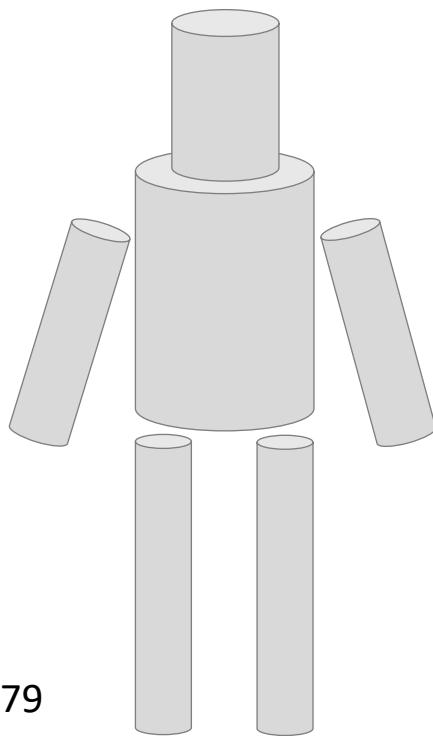
1963
Roberts

<https://dspace.mit.edu/handle/1721.1/6125>

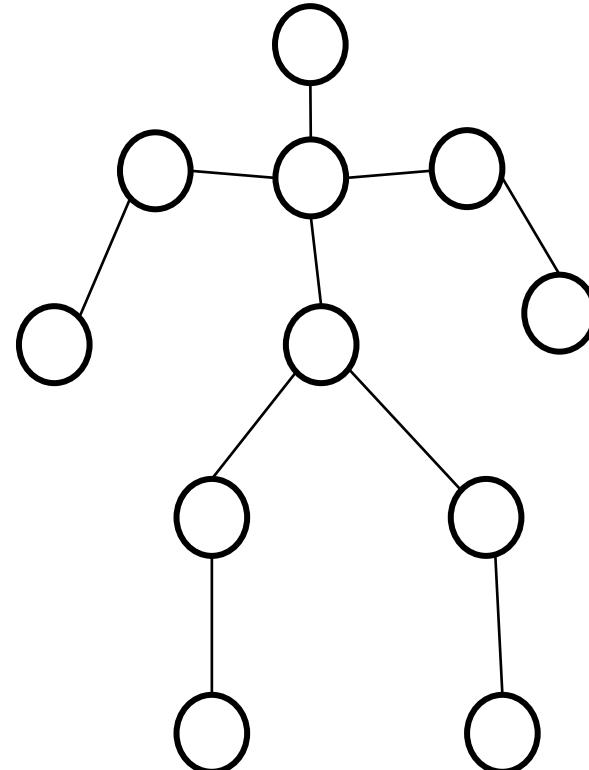


Stages of Visual Representation, David Marr, 1970s

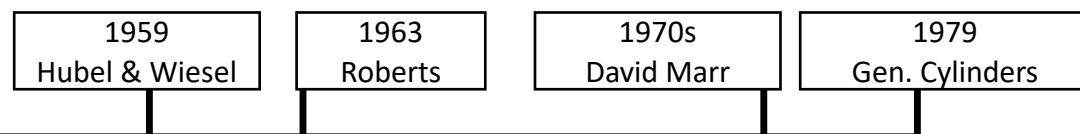
Recognition via Parts (1970s)



Generalized Cylinders,
Brooks and Binford, 1979



Pictorial Structures,
Fischler and Elshlager, 1973



Recognition via Edge Detection (1980s)



1959
Hubel & Wiesel

1963
Roberts

1970s
David Marr

1979
Gen. Cylinders

1986
Canny

John Canny, 1986
David Lowe, 1987

Image is CC0 1.0 public domain

Recognition via Grouping (1990s)



1959
Hubel & Wiesel

1963
Roberts

1970s
David Marr

1979
Gen. Cylinders

1986
Canny

1997
Norm. Cuts

AI Winter

Normalized Cuts, Shi and Malik, 1997

Recognition via Matching (2000s)



[Image](#) is public domain



[Image](#) is public domain

1959
Hubel & Wiesel

1963
Roberts

1970s
David Marr

1979
Gen. Cylinders

1986
Canny

1997
Norm. Cuts

1999
SIFT

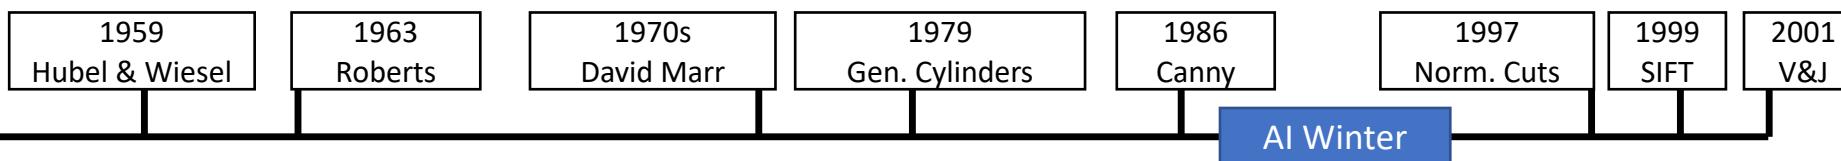
AI Winter

SIFT, David
Lowe, 1999

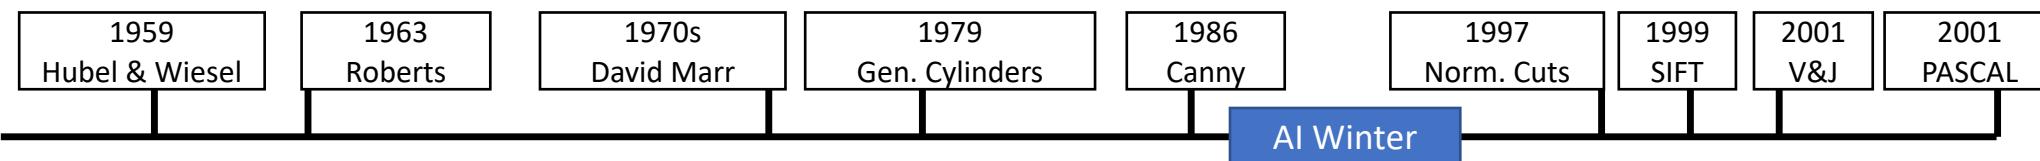
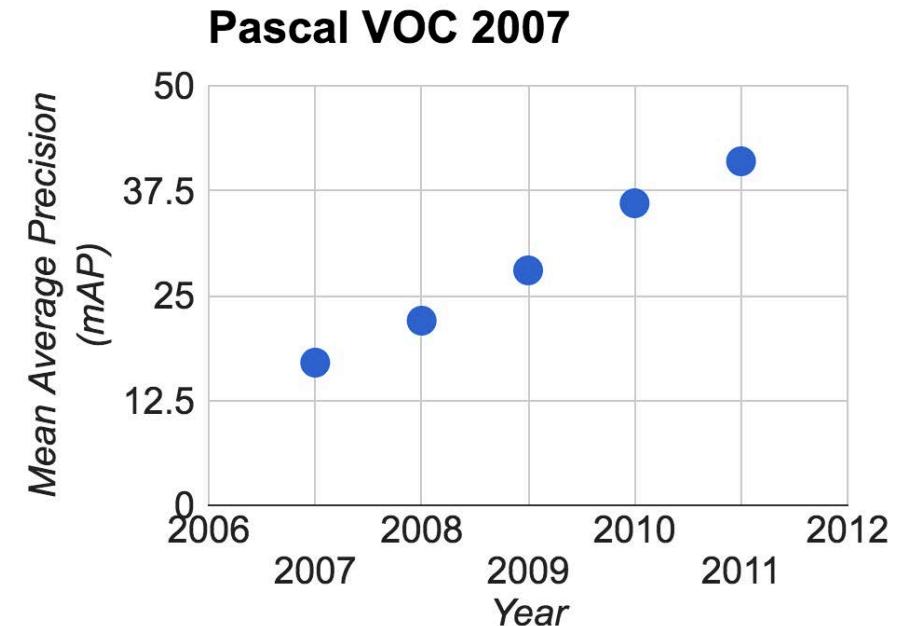
Face Detection

Viola and Jones, 2001

One of the first successful applications of machine learning to vision



PASCAL Visual Object Challenge



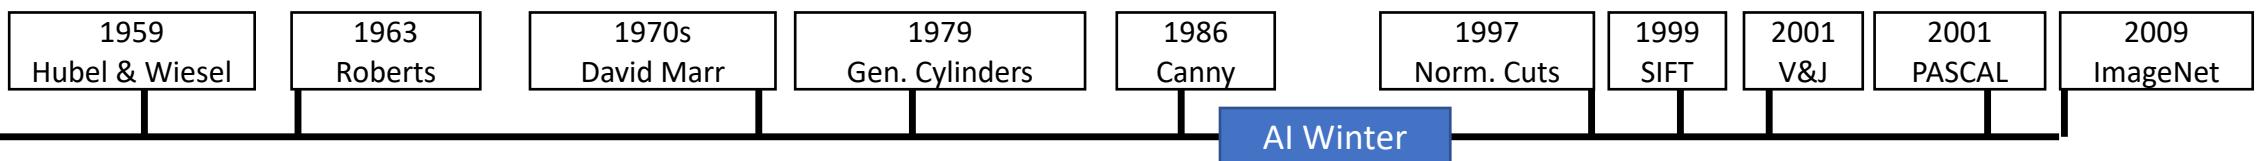
IMAGENET Large Scale Visual Recognition Challenge

The Image Classification Challenge:
1,000 object classes
1,431,167 images

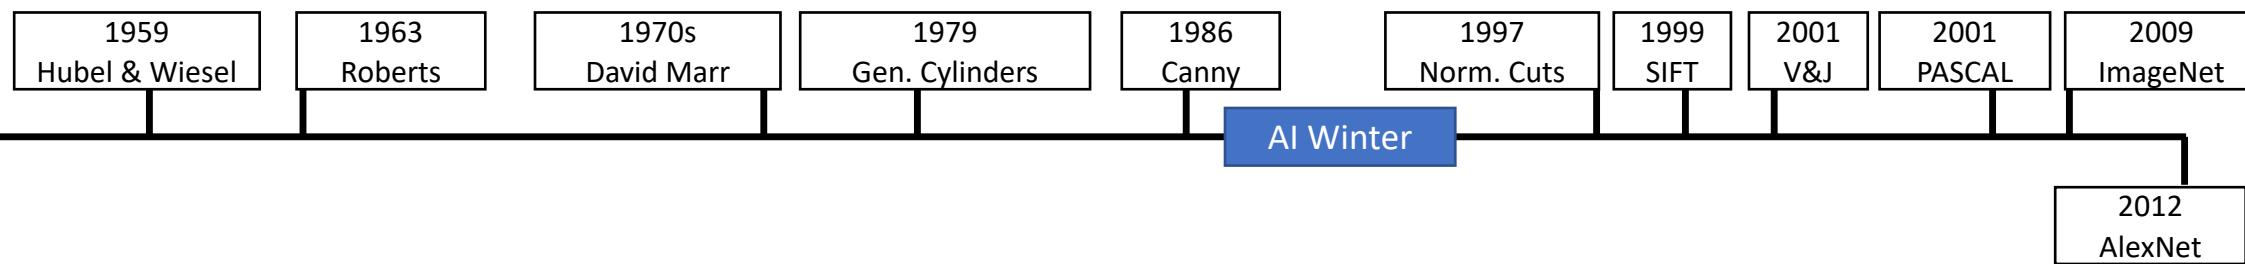
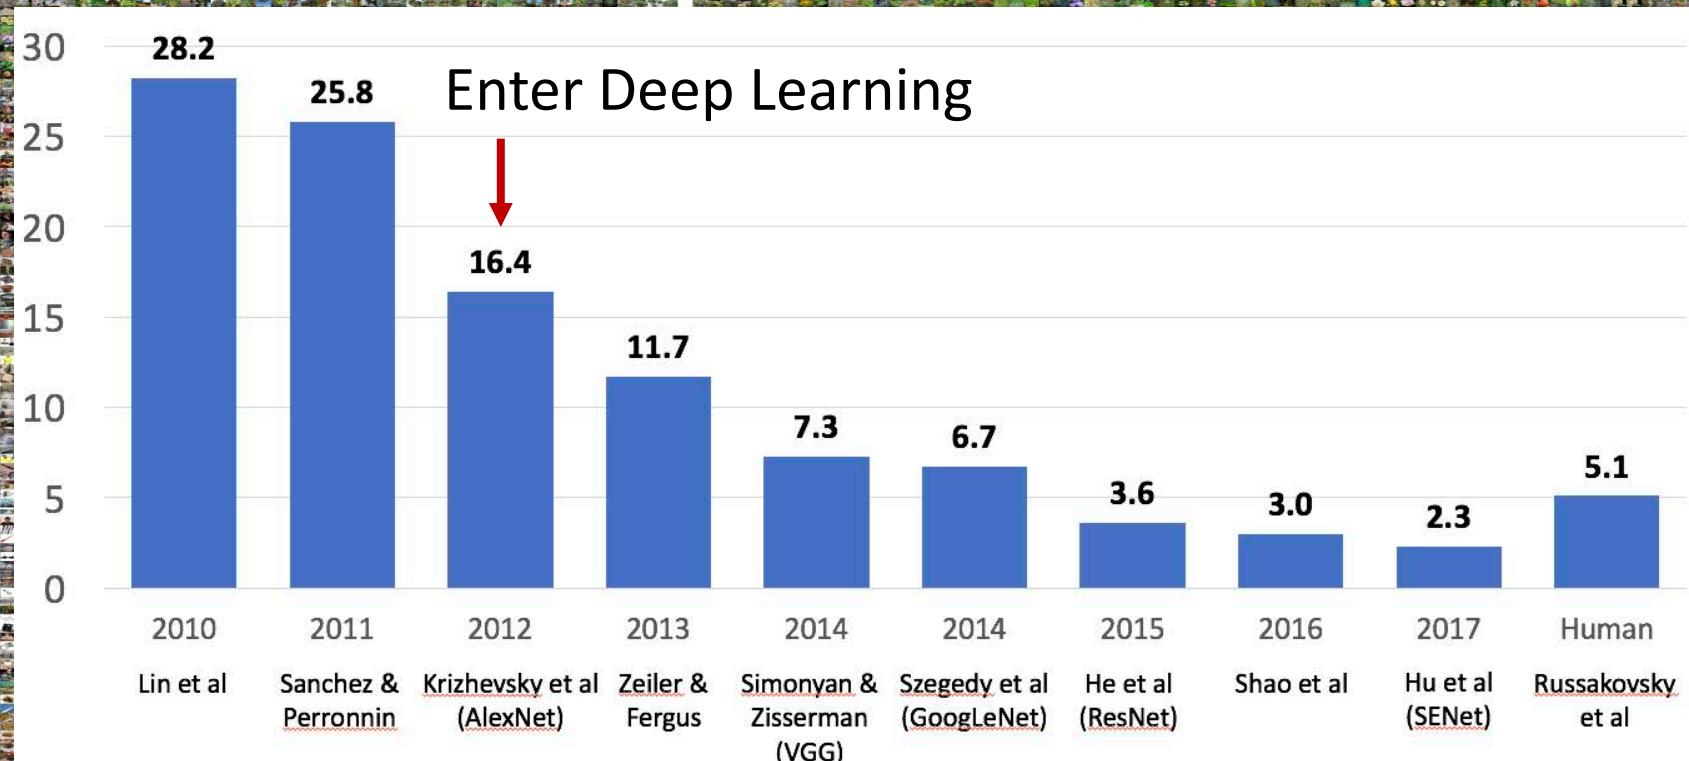


Output:
Scale
T-shirt
Steel drum
Drumstick
Mud turtle

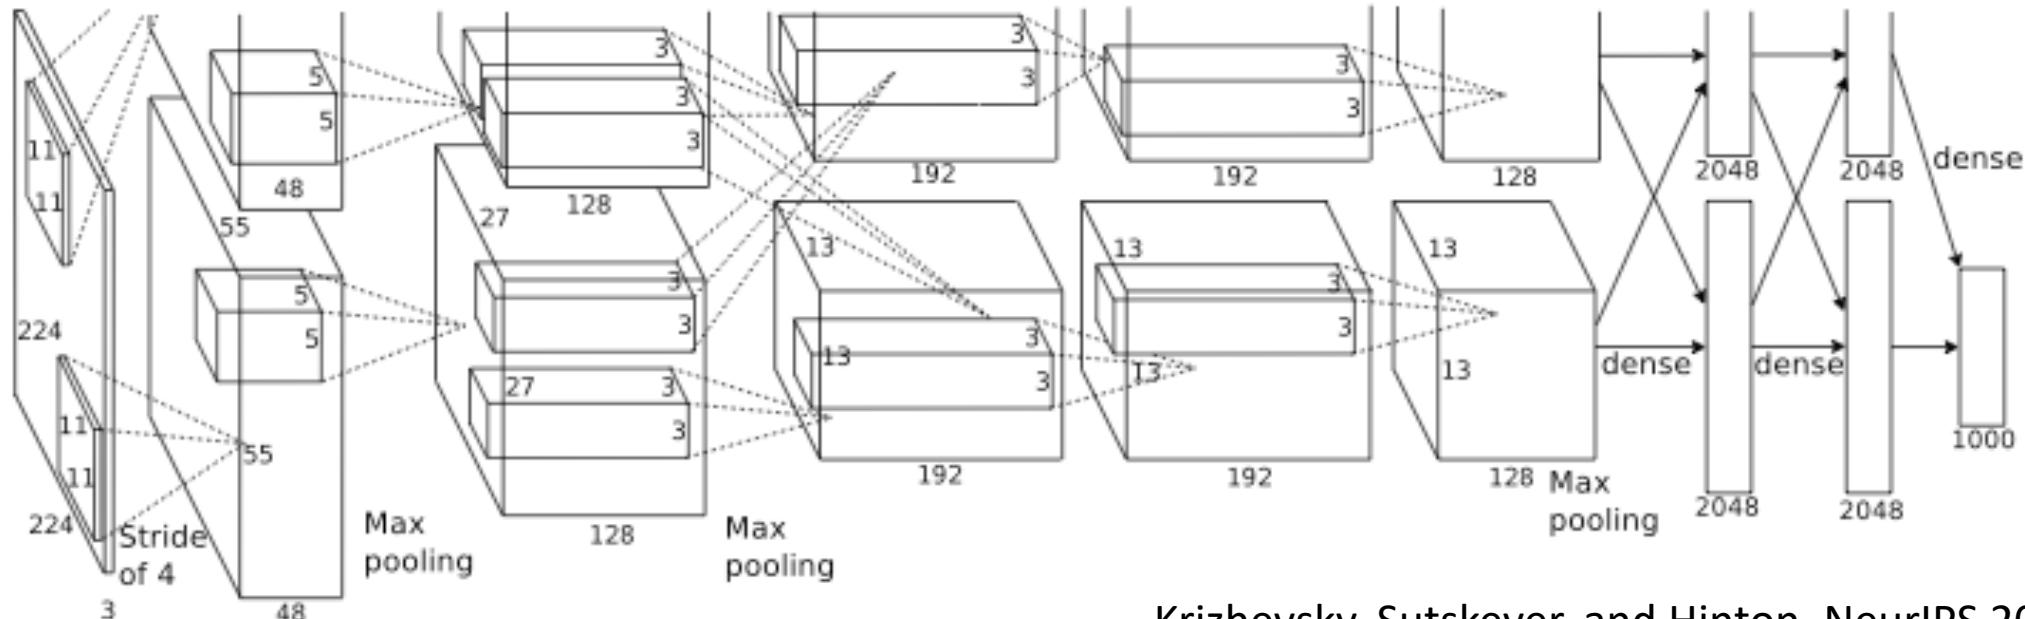
Deng et al, 2009
Russakovsky et al. IJCV 2015



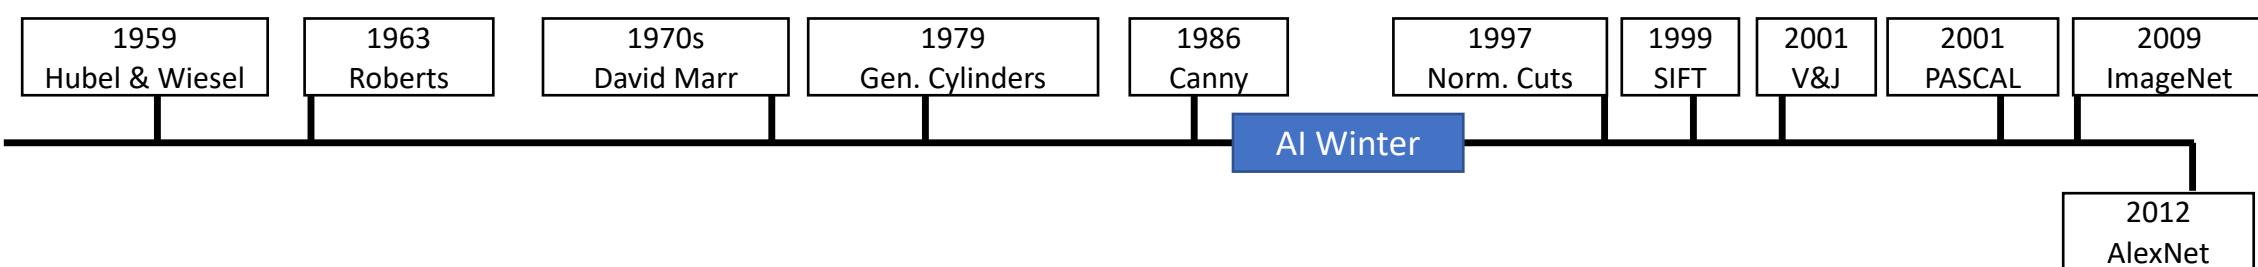
IMAGENET Large Scale Visual Recognition Challenge



AlexNet: Deep Learning Goes Mainstream



Krizhevsky, Sutskever, and Hinton, NeurIPS 2012



Perceptron

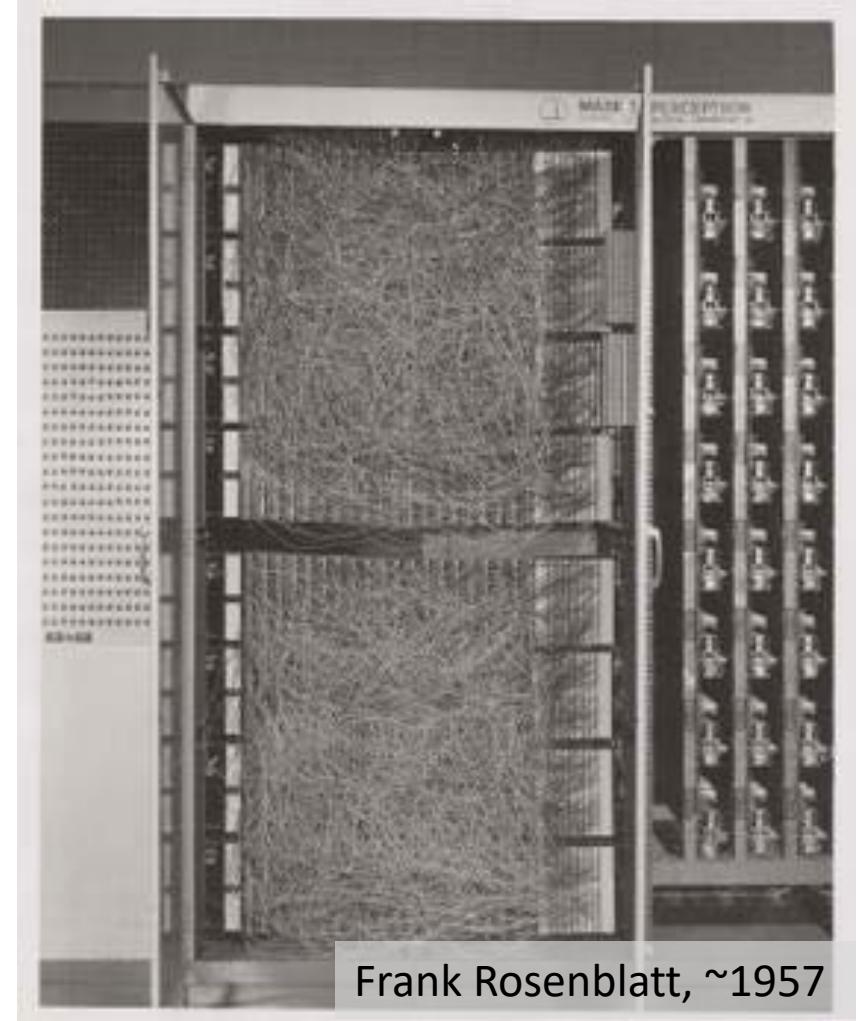
One of the earliest algorithms that could learn from data

Implemented in hardware! Weights stored in potentiometers,
updated with electric motors during learning

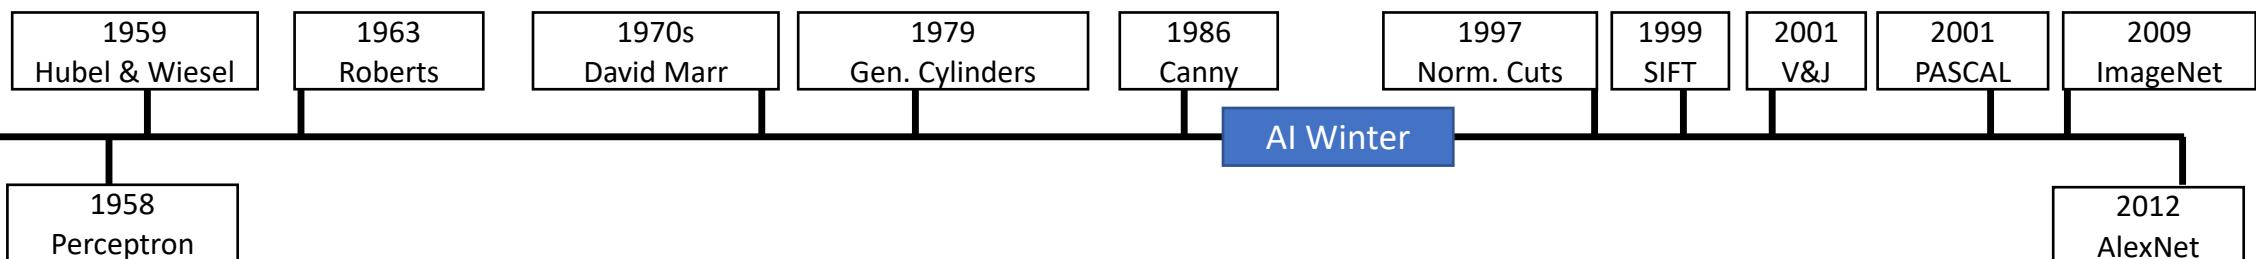
Connected to a camera that used 20x20 cadmium sulfide
photocells to make a 400-pixel image

Could learn to recognize letters of the alphabet

Today we would recognize it as a **linear classifier**

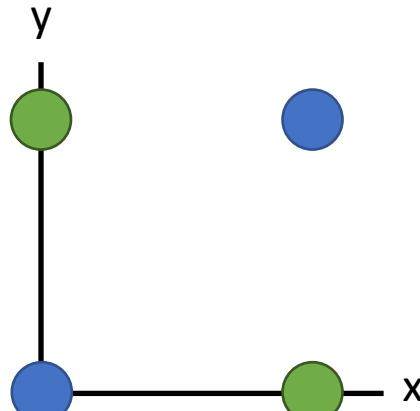


Frank Rosenblatt, ~1957

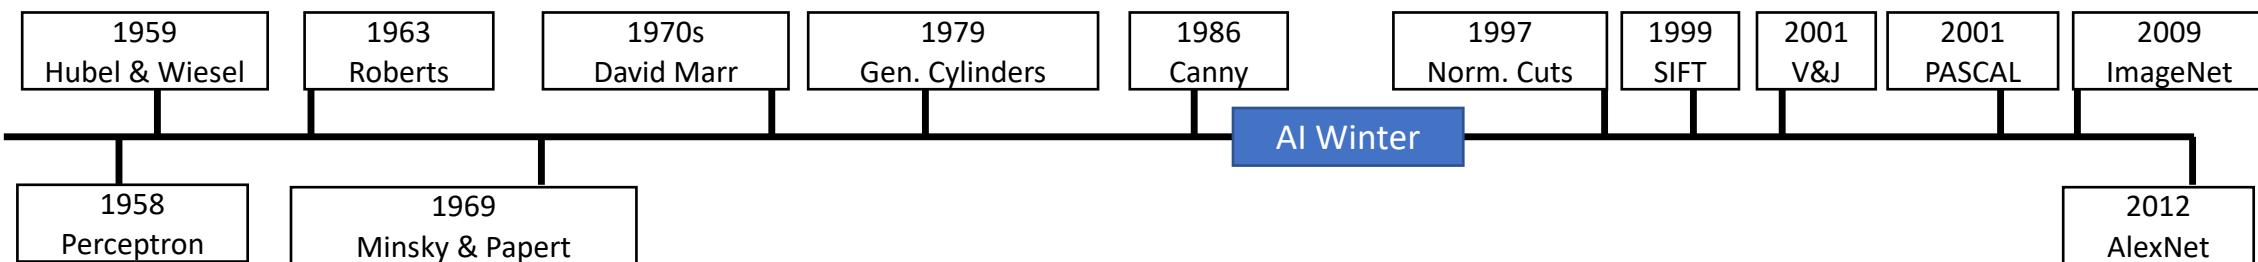
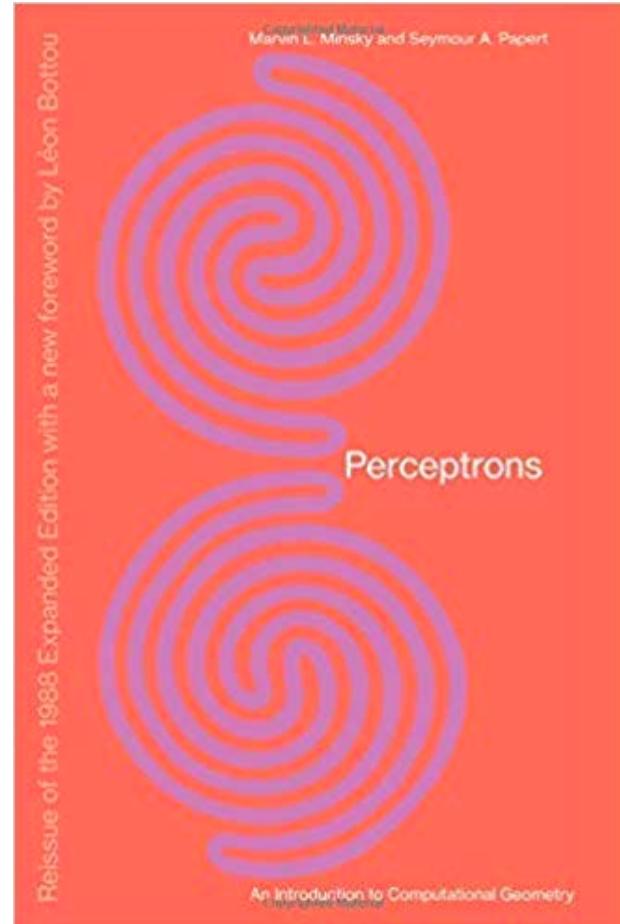


Minsky and Papert, 1969

X	Y	F(x,y)
0	0	0
0	1	1
1	0	1
1	1	0



Showed that Perceptrons could not learn the XOR function
Caused a lot of disillusionment in the field

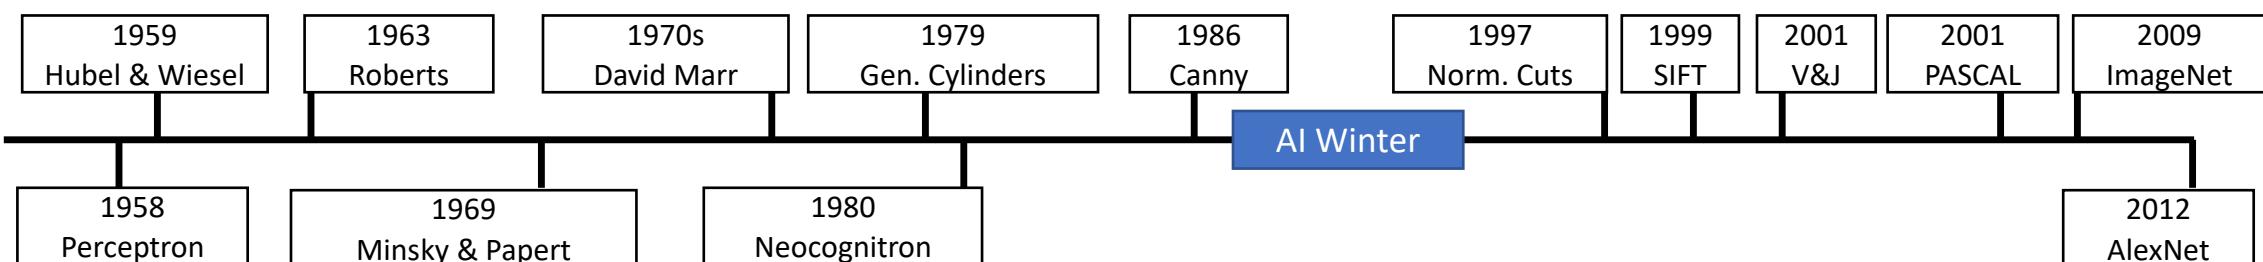
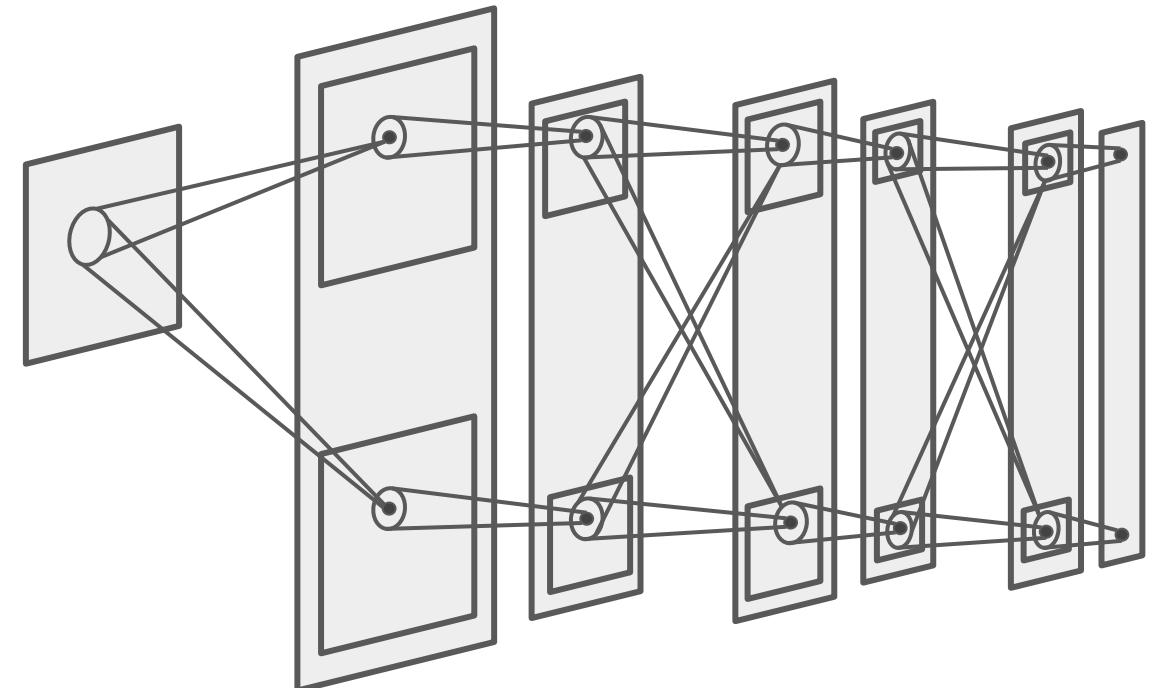


Neocognitron: Fukushima, 1980

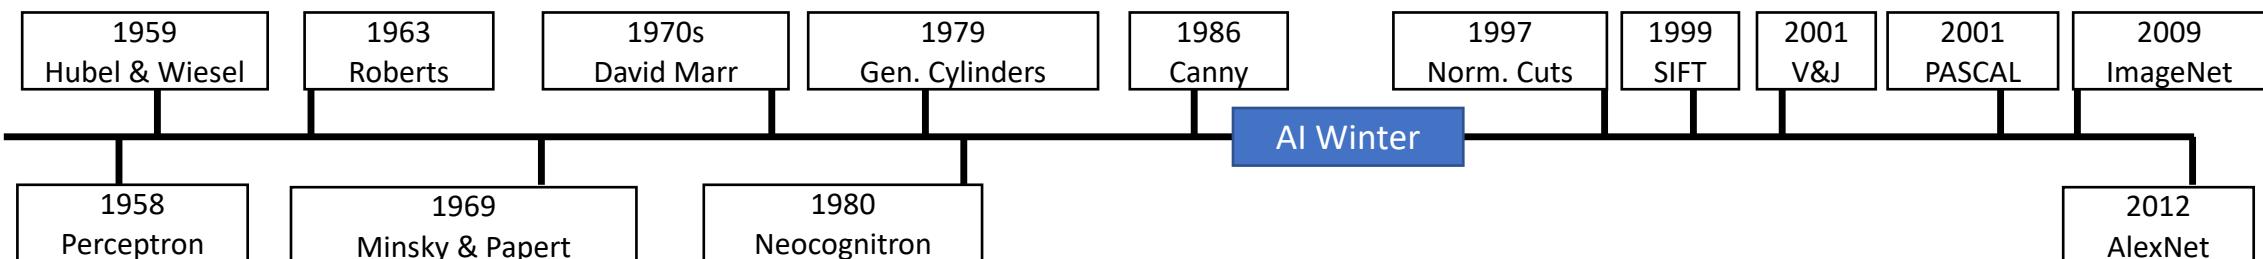
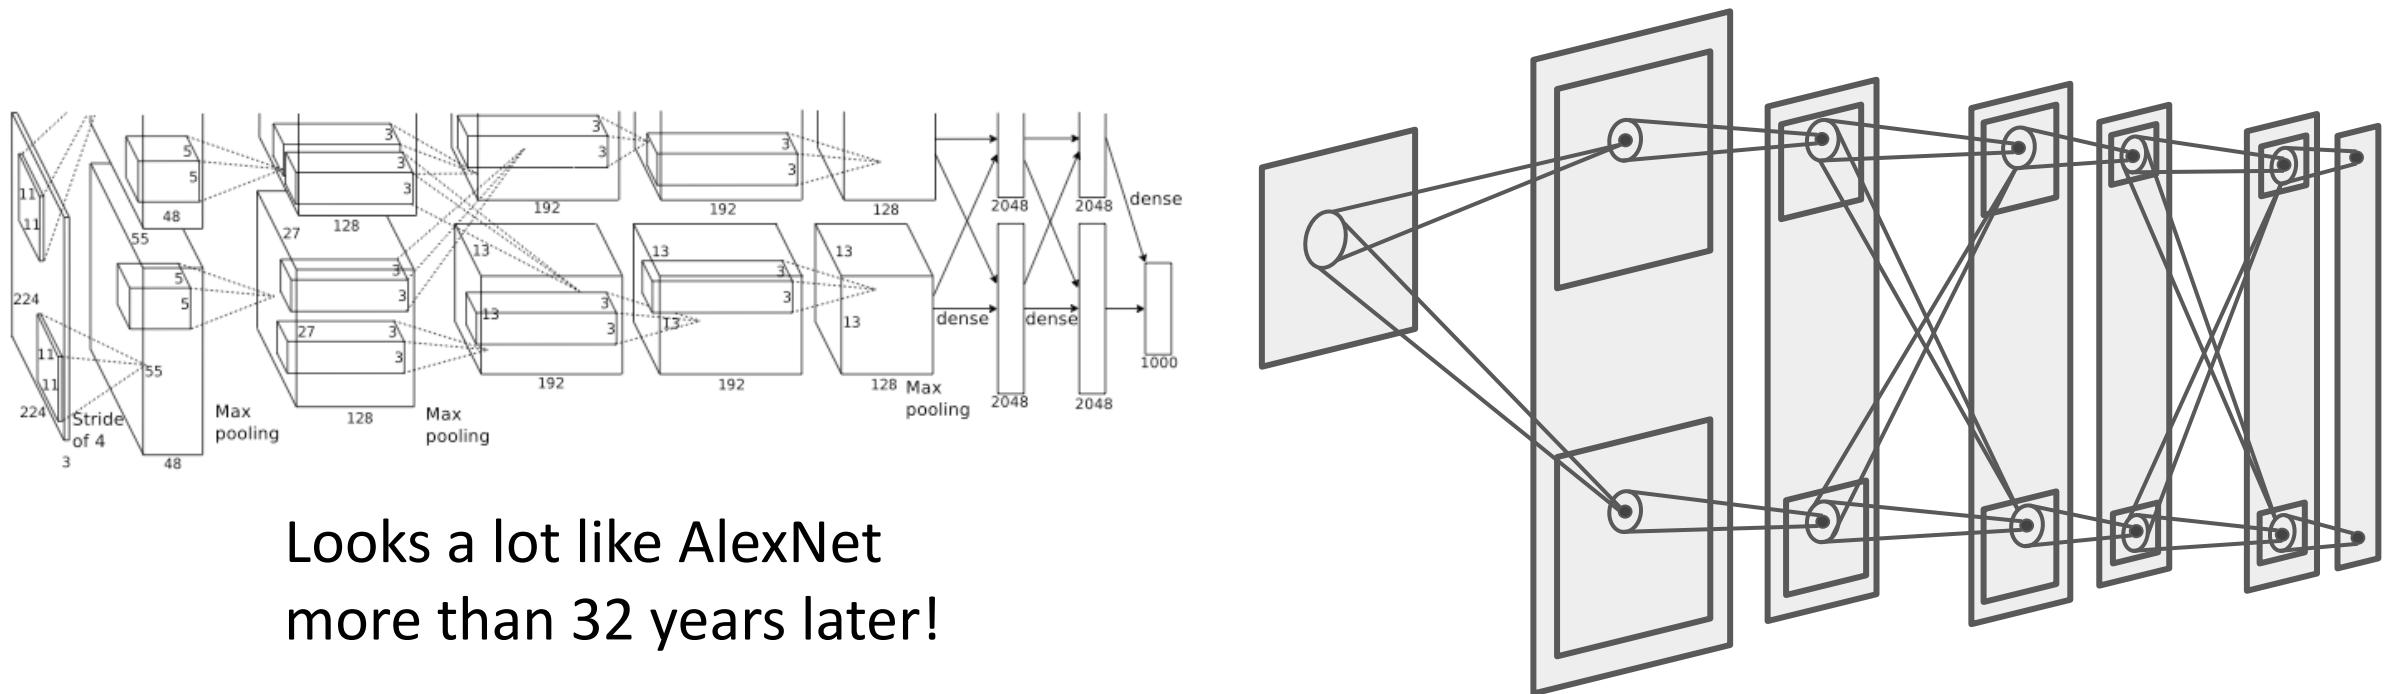
Computational model the visual system,
directly inspired by Hubel and Wiesel's
hierarchy of complex and simple cells

Interleaved simple cells (convolution)
and complex cells (pooling)

No practical training algorithm



Neocognitron: Fukushima, 1980



Backprop: Rumelhart, Hinton, and Williams, 1986

Introduced backpropagation
for computing gradients in
neural networks

Successfully trained
perceptrons with multiple
layers

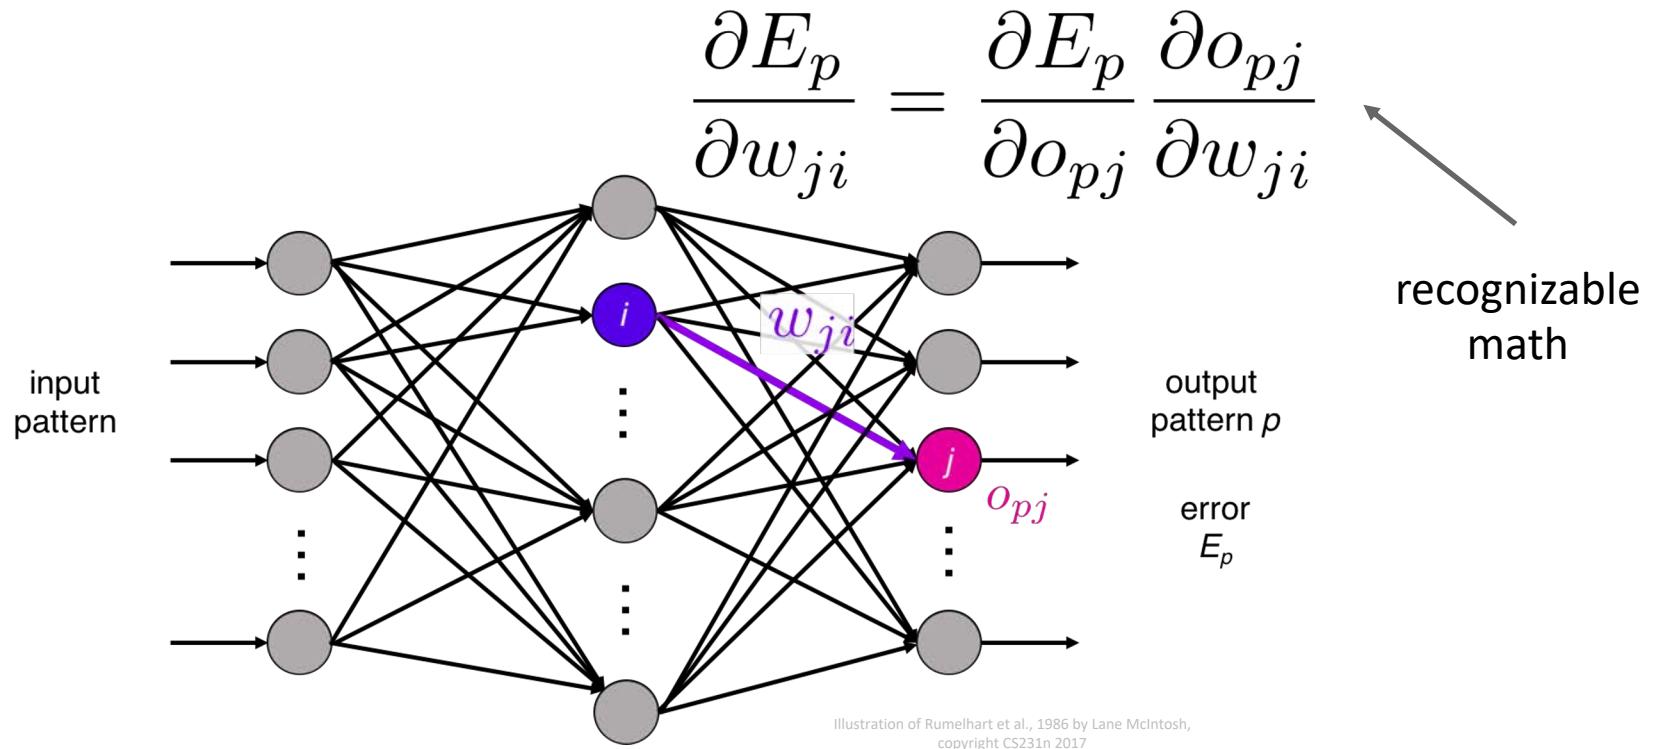
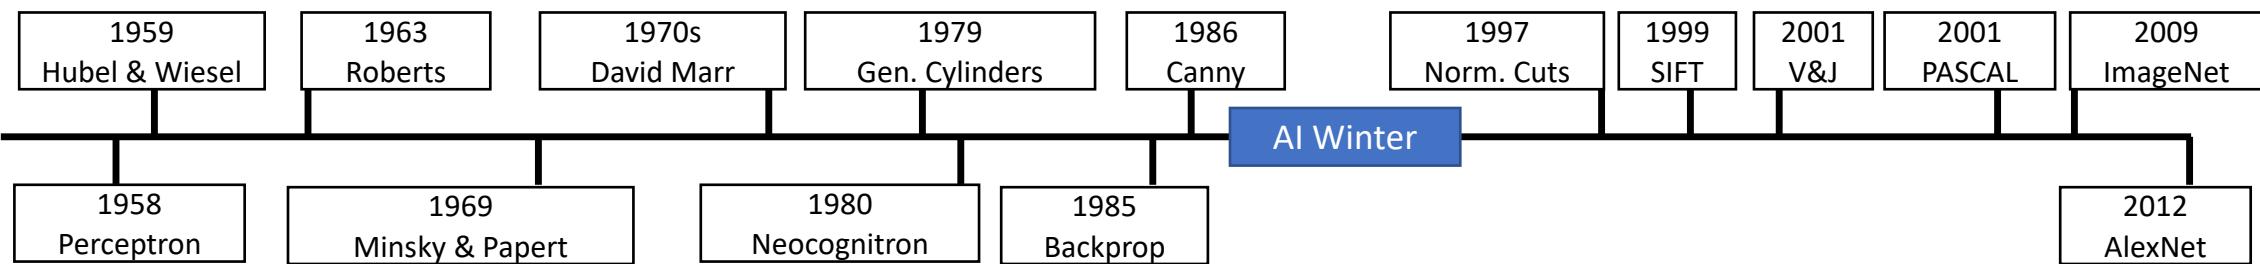
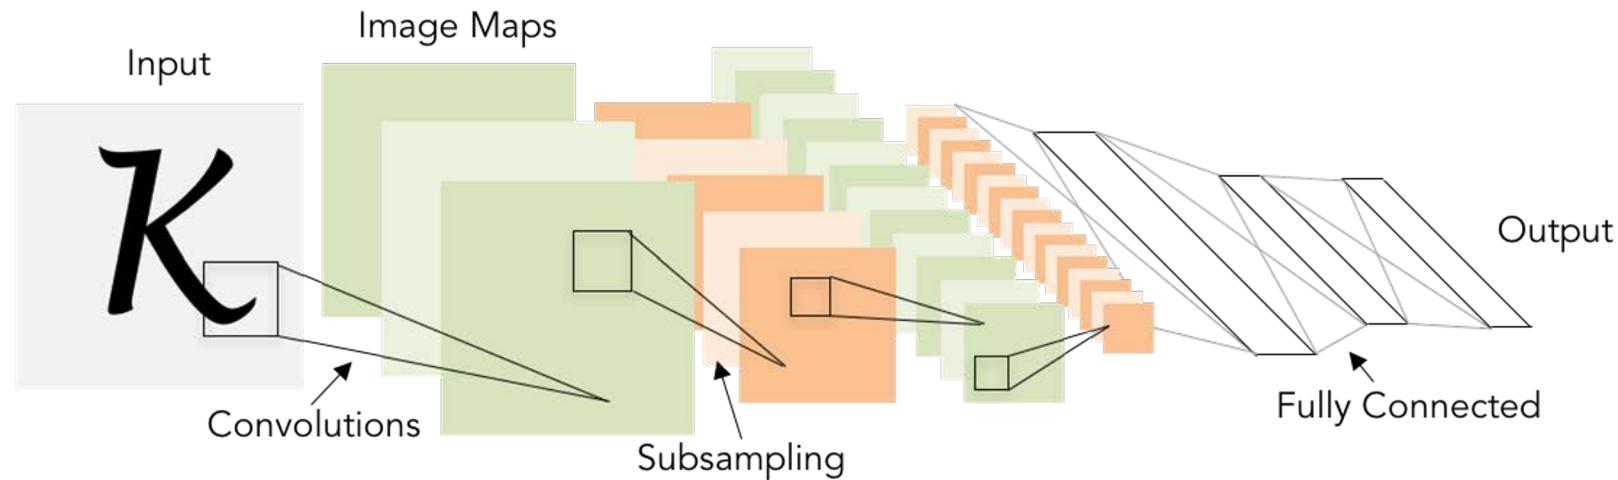


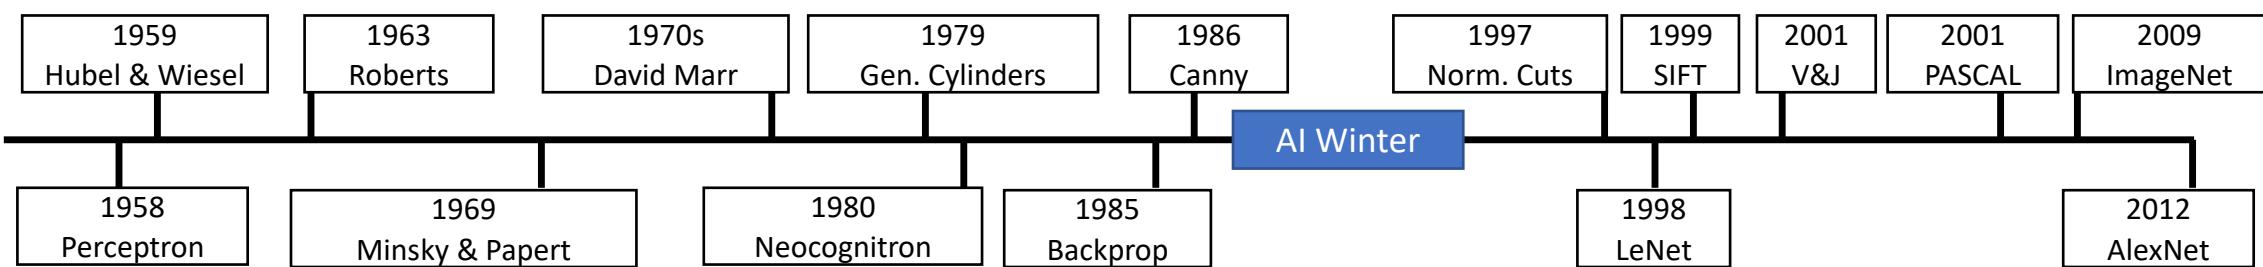
Illustration of Rumelhart et al., 1986 by Lane McIntosh,
copyright CS231n 2017



Convolutional Networks: LeCun et al, 1998



Applied backprop algorithm to a Neocognitron-like architecture
Learned to recognize handwritten digits
Was deployed in a commercial system by NEC, processed handwritten checks
Very similar to our modern convolutional networks!

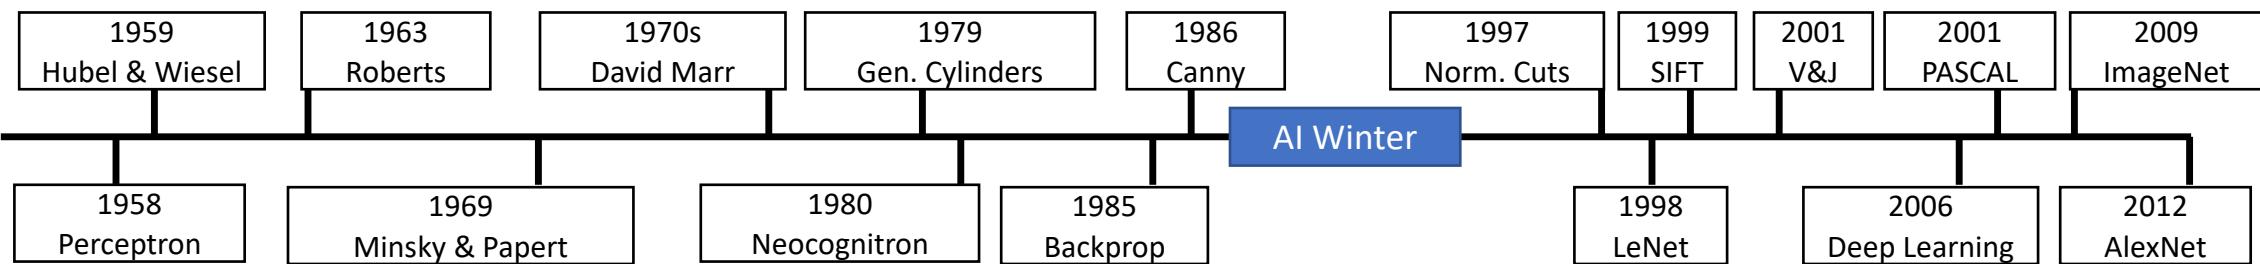
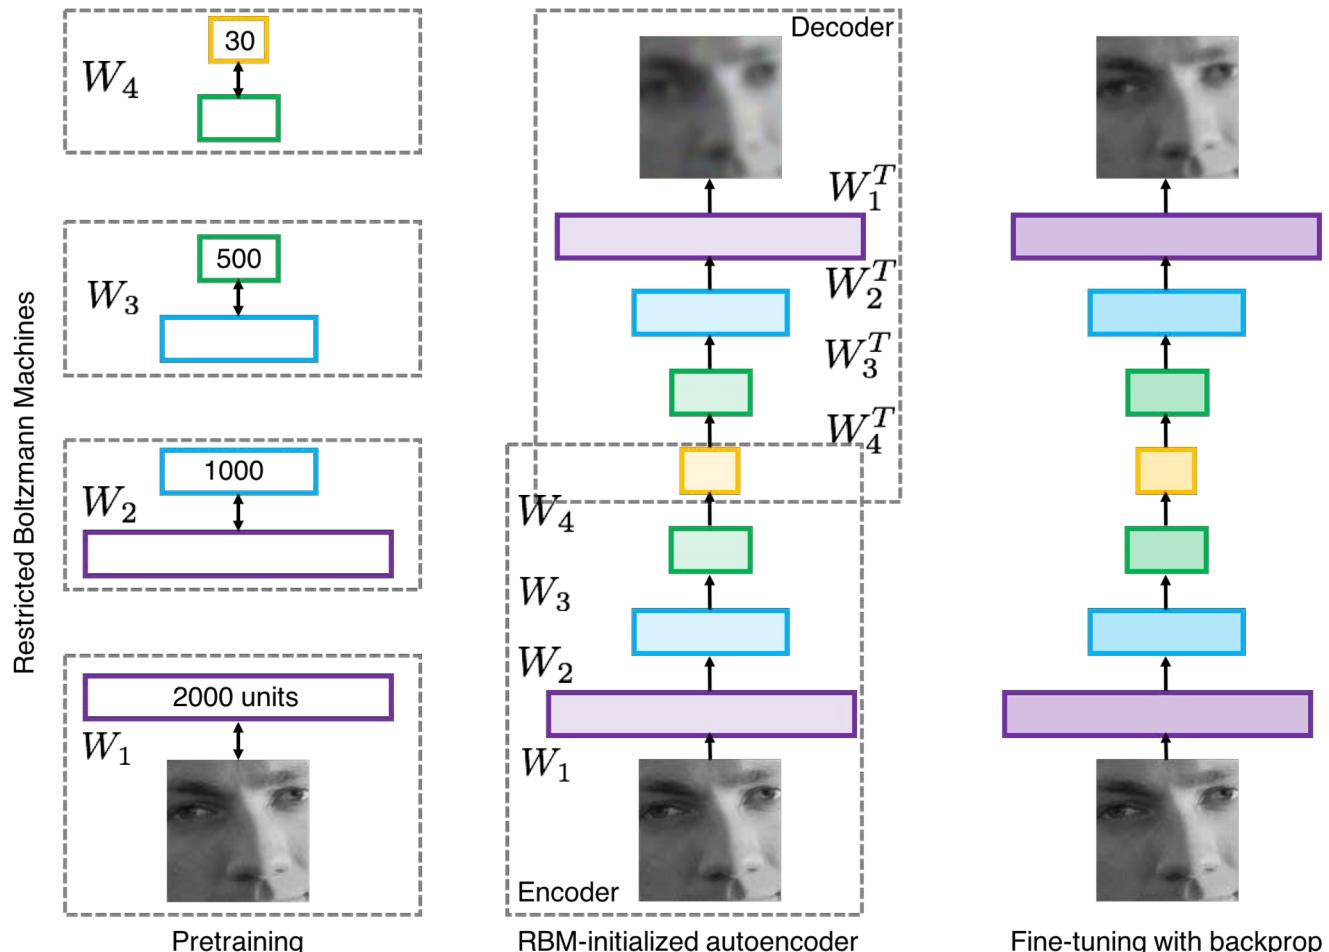


2000s: “Deep Learning”

People tried to train neural networks that were deeper and deeper

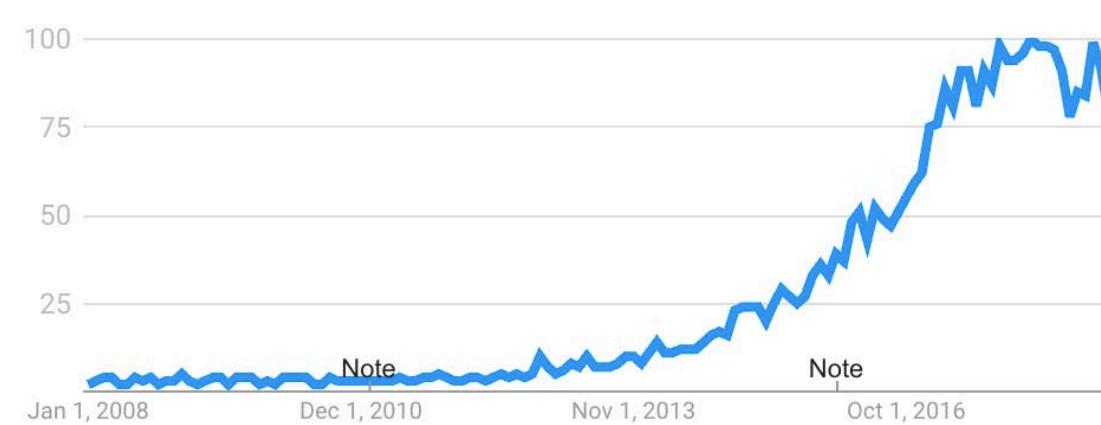
Not a mainstream research topic at this time

Hinton and Salakhutdinov, 2006
Bengio et al, 2007
Lee et al, 2009
Glorot and Bengio, 2010



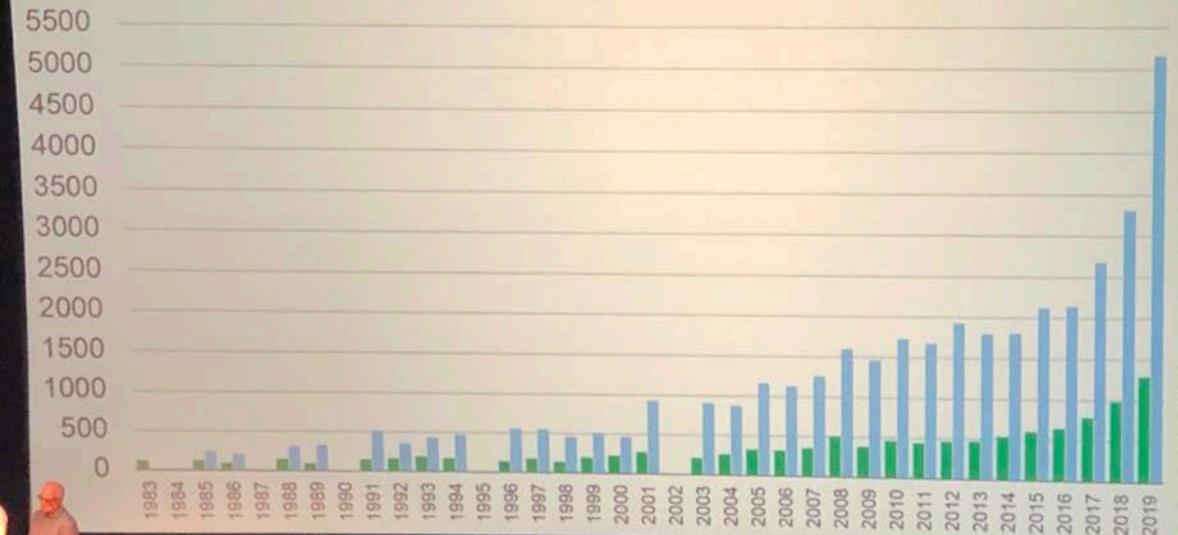
2012 to Present: Deep Learning Explosion

Interest over time

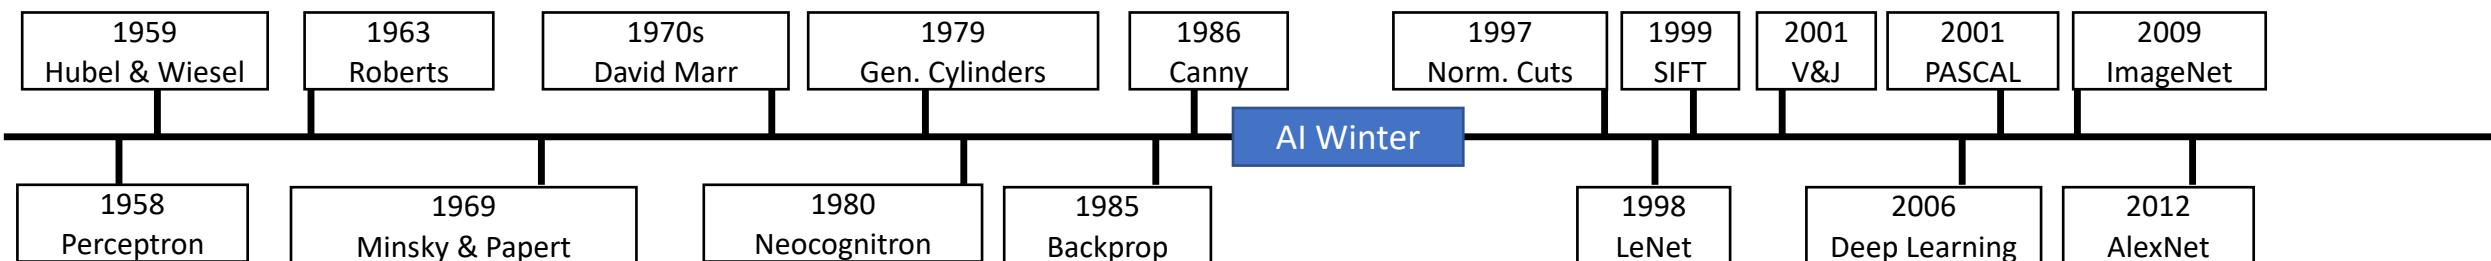


Google Trends: “Deep Learning”

CVPR Submitted and Accepted Papers



Publications at top Computer Vision conference



2012 to Present: ConvNets are everywhere

Image Classification



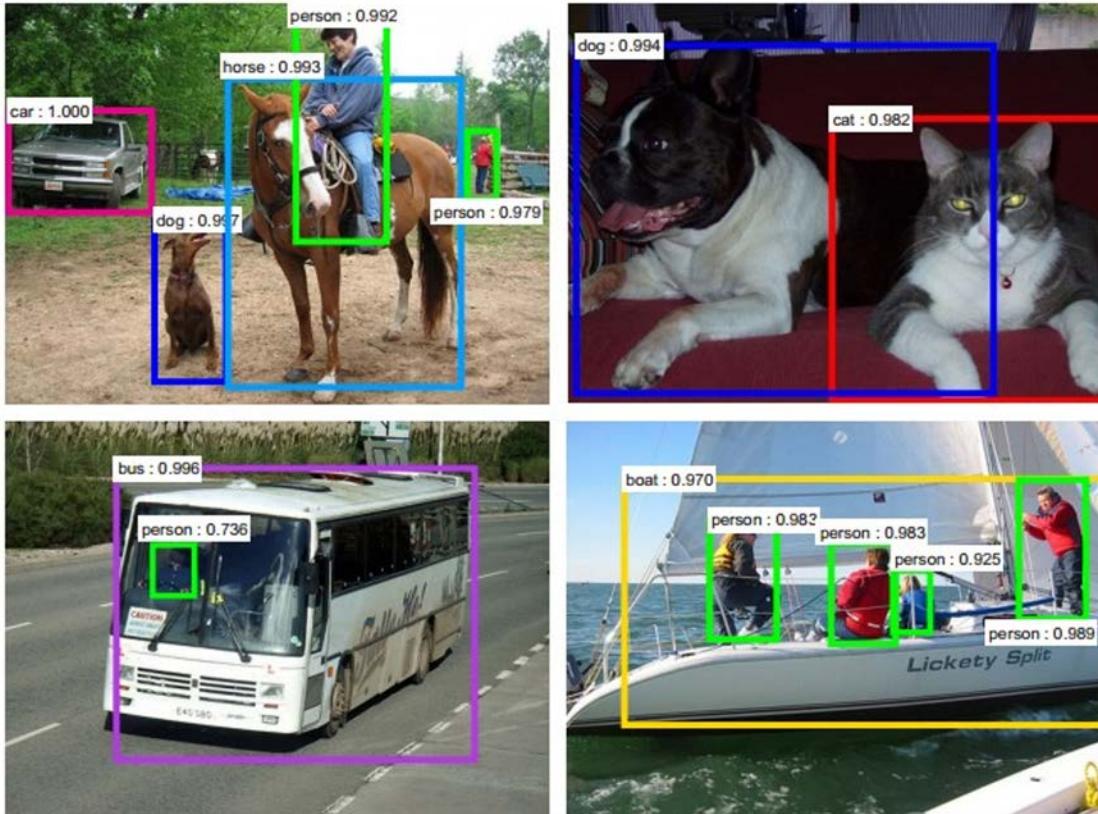
Image Retrieval



Figures copyright Alex Krizhevsky, Ilya Sutskever, and Geoffrey Hinton, 2012. Reproduced with permission.

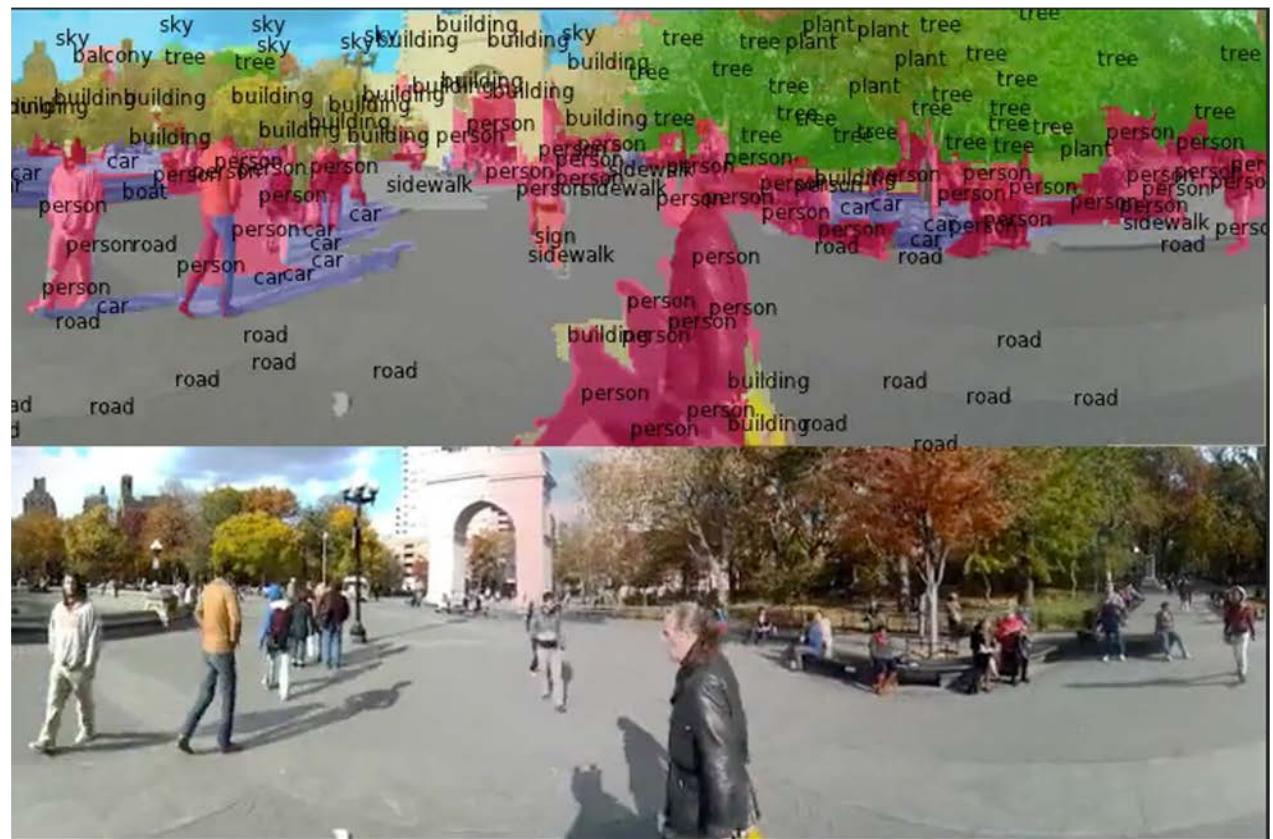
2012 to Present: ConvNets are everywhere

Object Detection



Ren, He, Girshick, and Sun, 2015

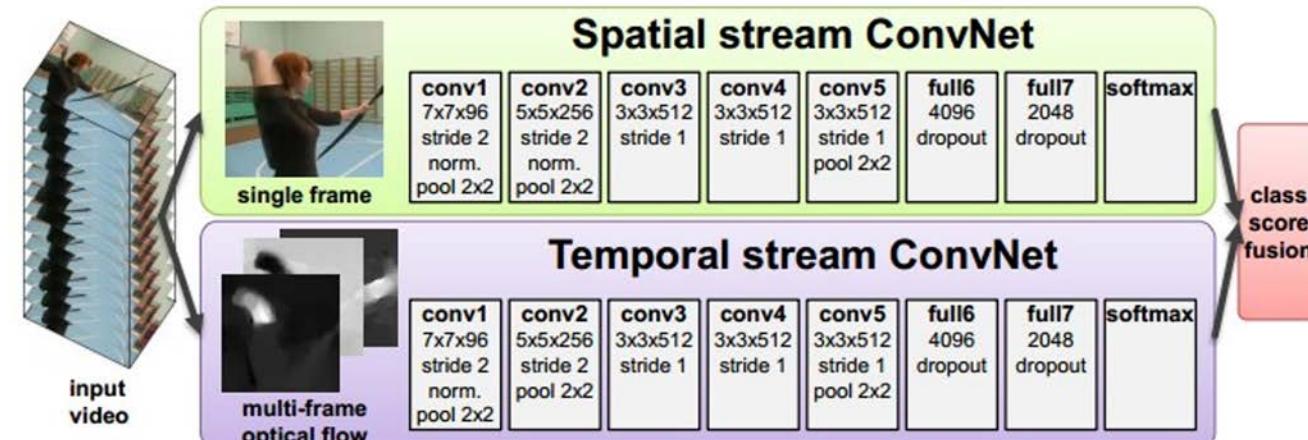
Image Segmentation



Fabaret et al, 2012

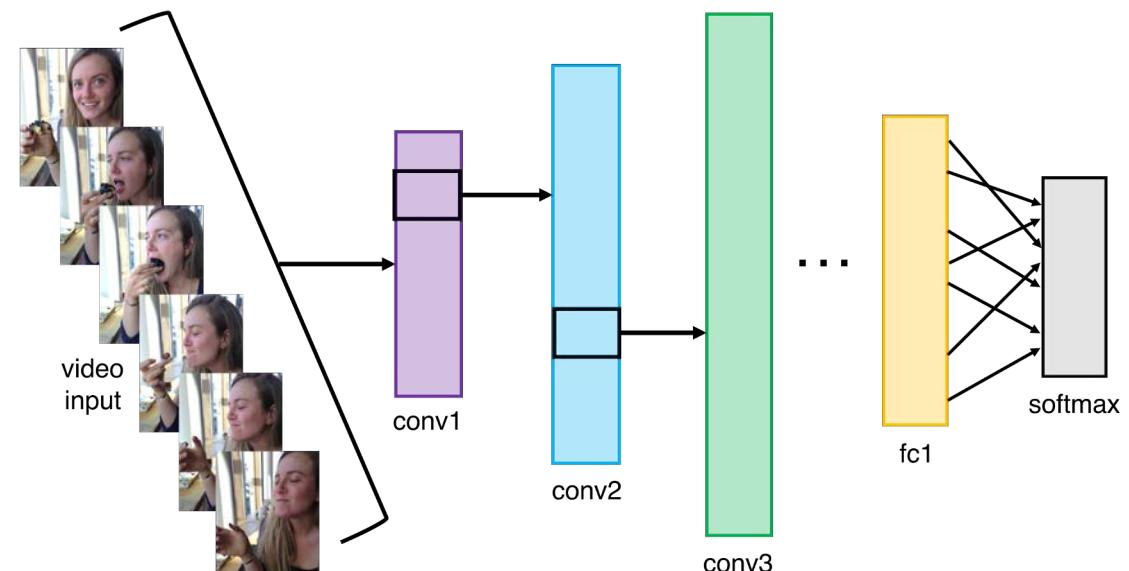
2012 to Present: ConvNets are everywhere

Video Classification



Simonyan et al, 2014

Activity Recognition

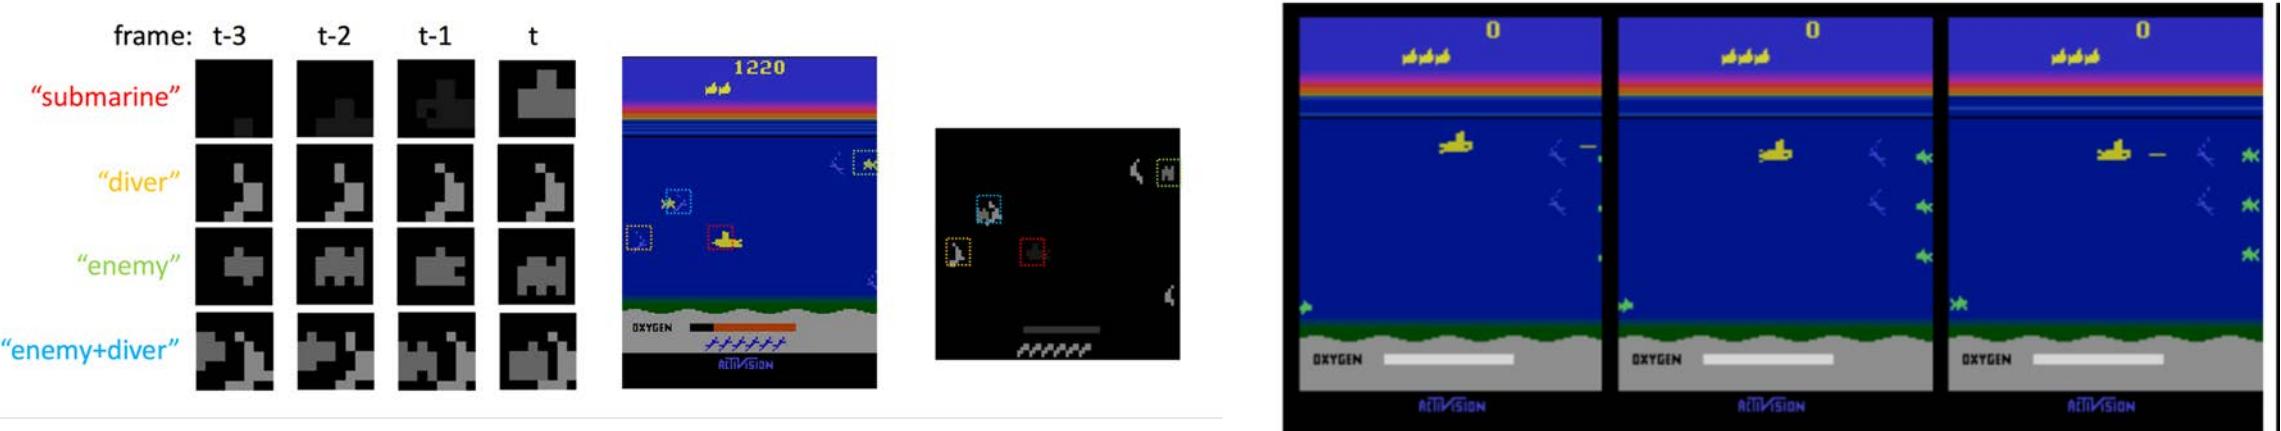


2012 to Present: ConvNets are everywhere

Pose Recognition (Toshev and Szegedy, 2014)

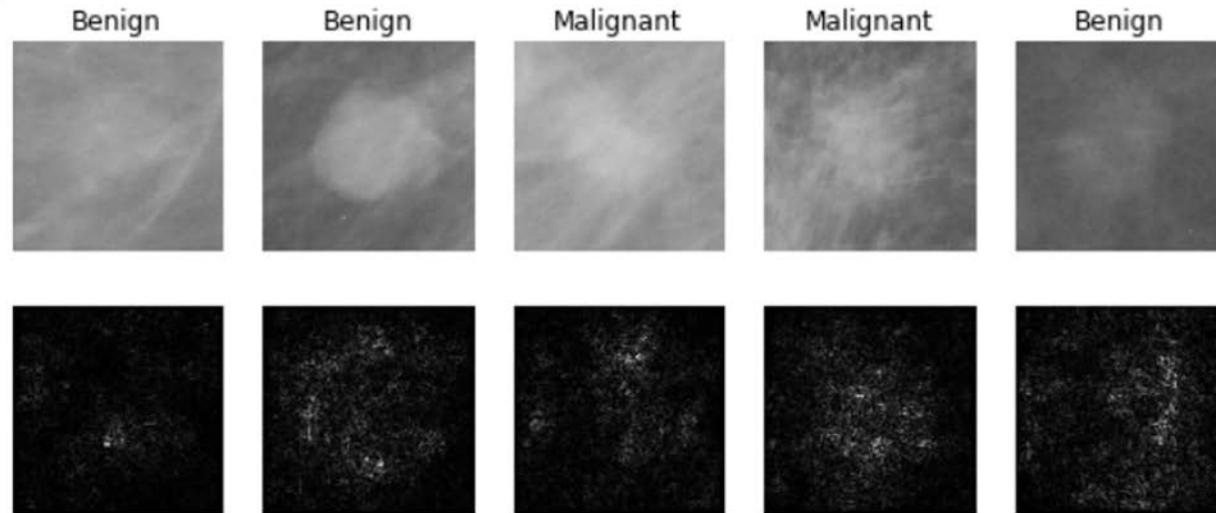


Playing Atari games (Guo et al, 2014)



2012 to Present: ConvNets are everywhere

Medical Imaging



Levy et al, 2016

Figure reproduced with permission

Whale recognition



[Kaggle Challenge](#)

This image by Christin Khan is in the public domain and originally came from the U.S. NOAA.

Dieleman et al, 2014

From left to right: public domain by NASA, usage permitted by
ESA/Hubble, public domain by NASA, and public domain.



2012 to Present: ConvNets are everywhere



*A white teddy bear
sitting in the grass*



*A man in a baseball
uniform throwing a ball*



*A woman is holding
a cat in her hand*



*A man riding a wave
on top of a surfboard*



*A cat sitting on a
suitcase on the floor*



*A woman standing on a
beach holding a surfboard*

Image Captioning

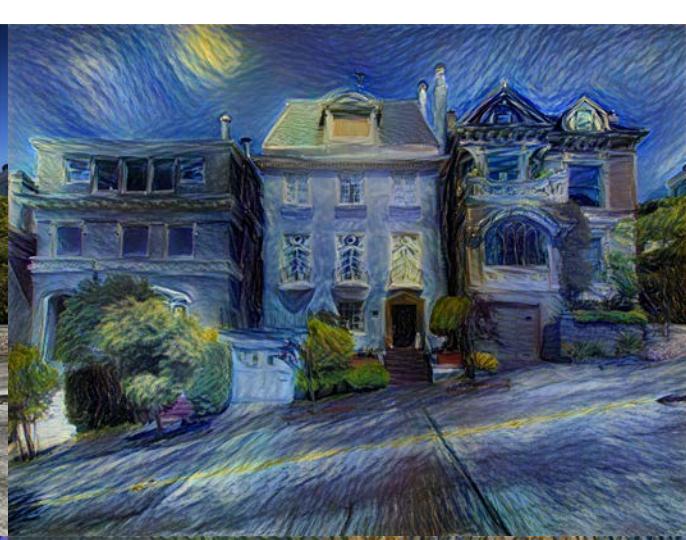
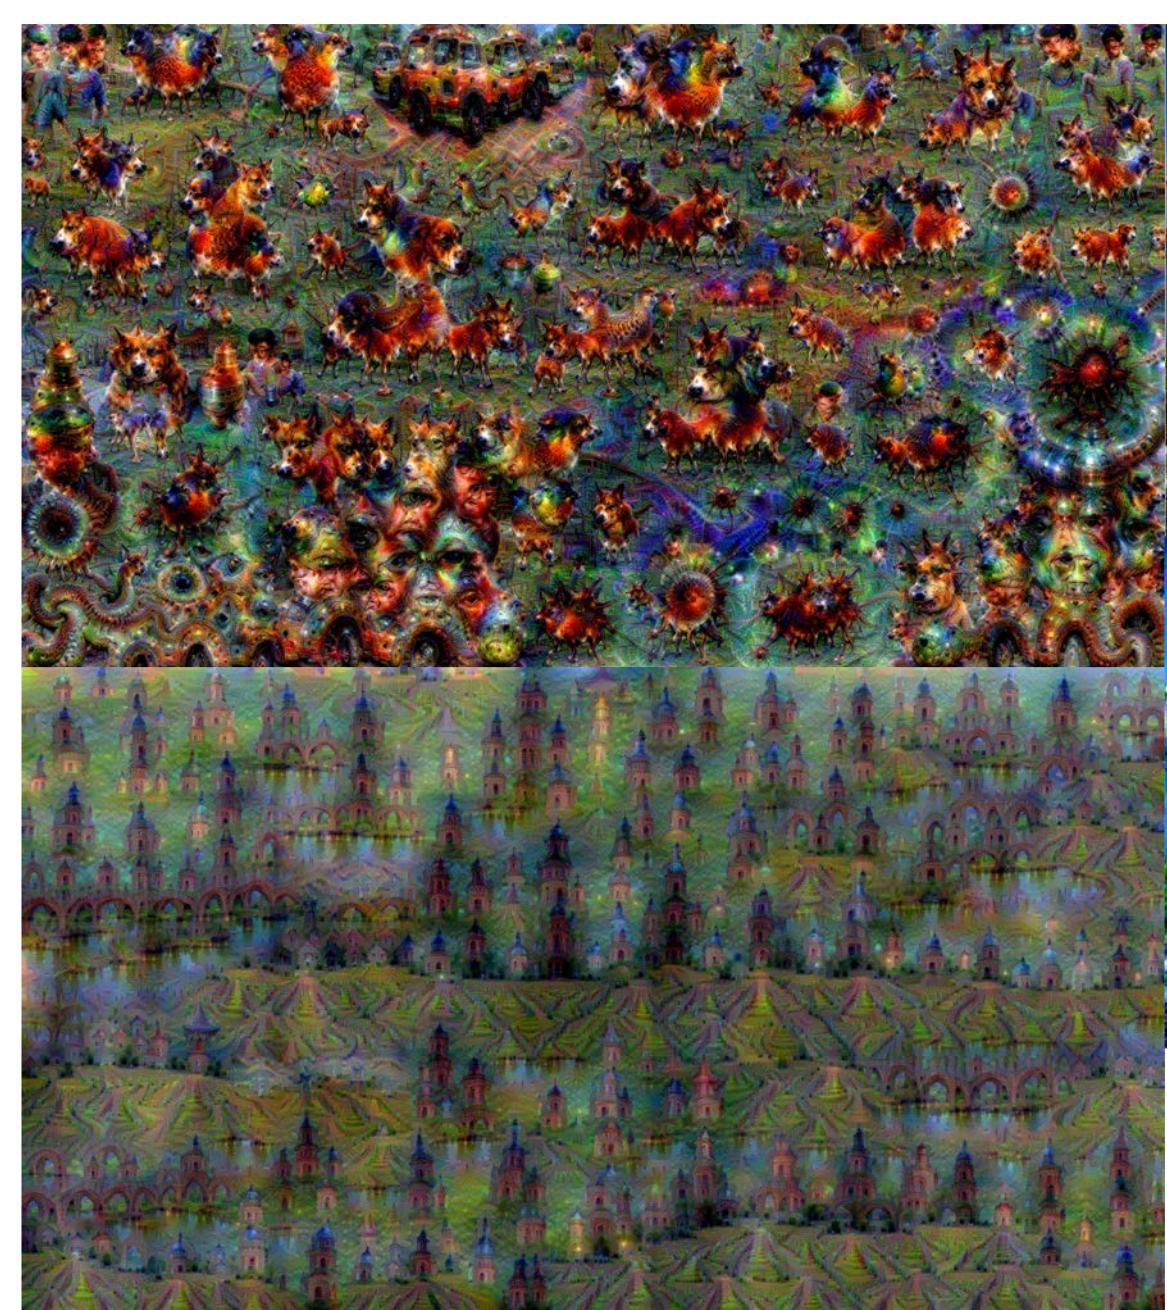
Vinyals et al, 2015

Karpathy and Fei-Fei, 2015

All images are CC0 Public domain:

<https://pixabay.com/en/luggage-antique-cat-1643010/>
<https://pixabay.com/en/teddy-plush-bears-cute-teddy-bear-1623436/>
<https://pixabay.com/en/surf-wave-summer-sport-litoral-1668716/>
<https://pixabay.com/en/woman-female-model-portrait-adult-983967/>
<https://pixabay.com/en/handstand-lake-meditation-496008/>
<https://pixabay.com/en/baseball-player-shortstop-infield-1045263/>

Captions generated by Justin Johnson using Neuraltalk2



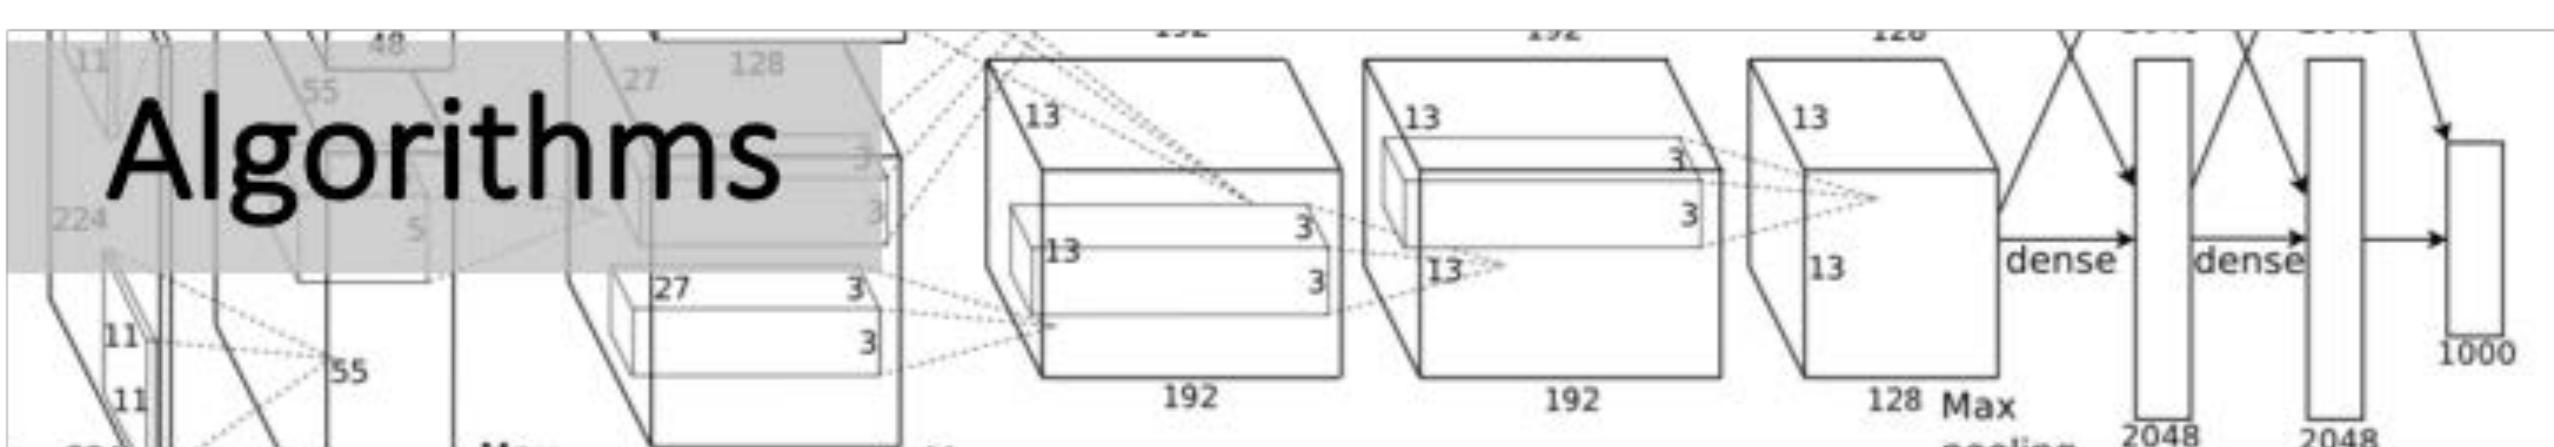
[Original image](#) is CC0 public domain
[Starry Night](#) and [Tree Roots](#) by Van Gogh are in the public domain
[Bokeh image](#) is in the public domain
Stylized images copyright Justin Johnson, 2017;
reproduced with permission



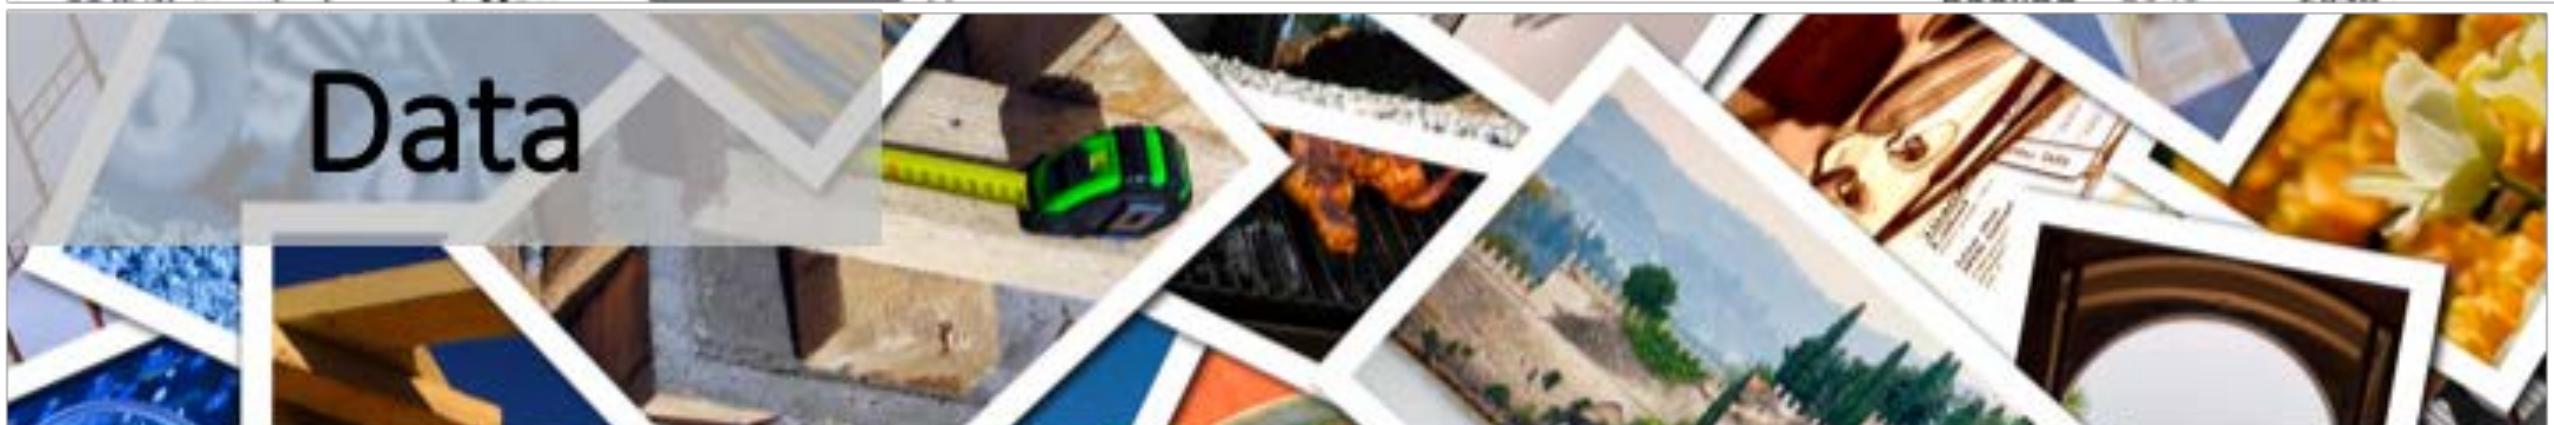
Mordvinsev et al, 2015
Gatys et al, 2016

Figures copyright Justin Johnson, 2015. Reproduced with permission. Generated using the Inceptionism approach from a [blog post](#) by Google Research.

Algorithms



Data

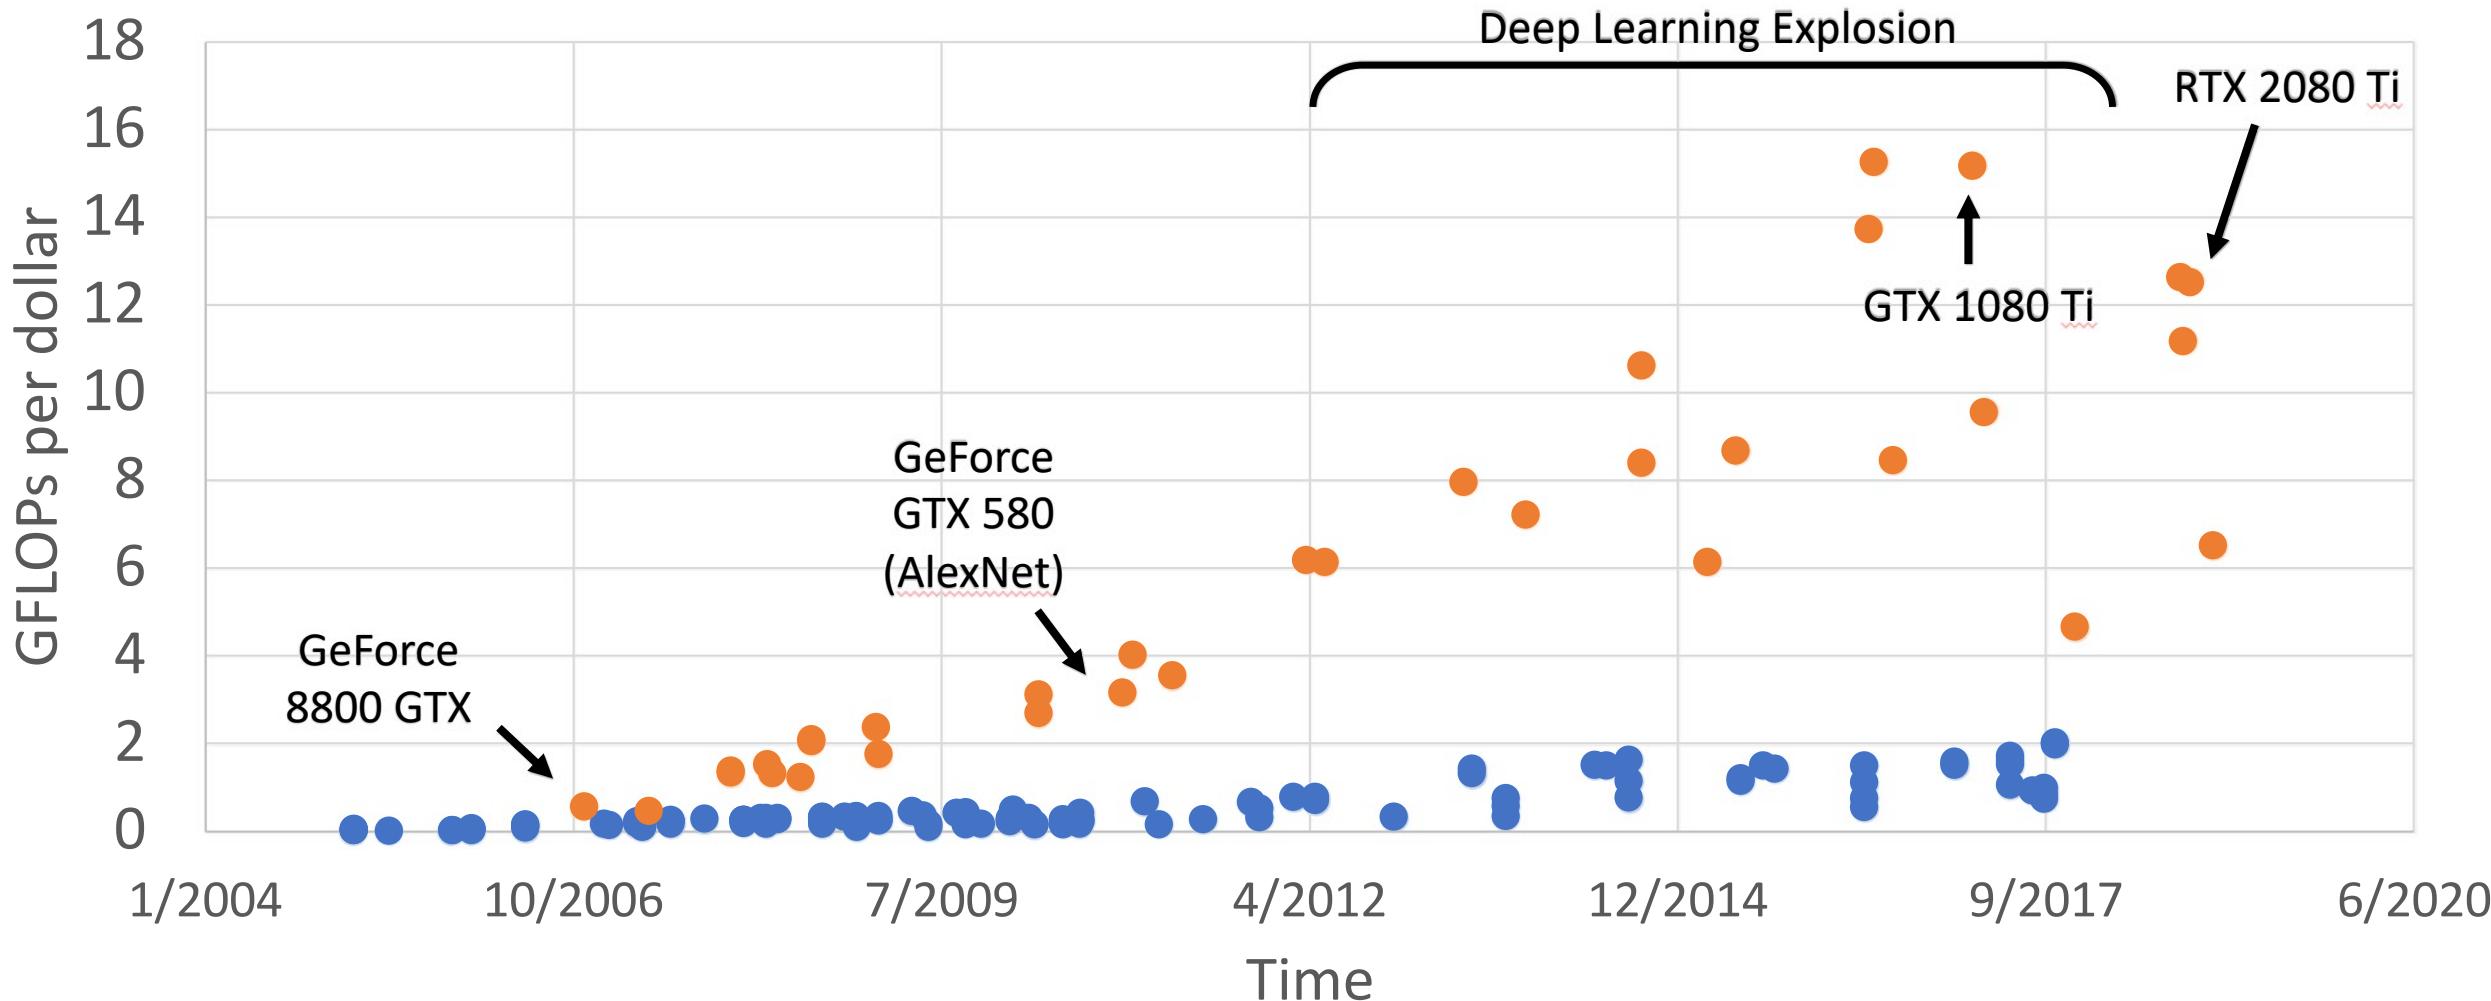


Computation



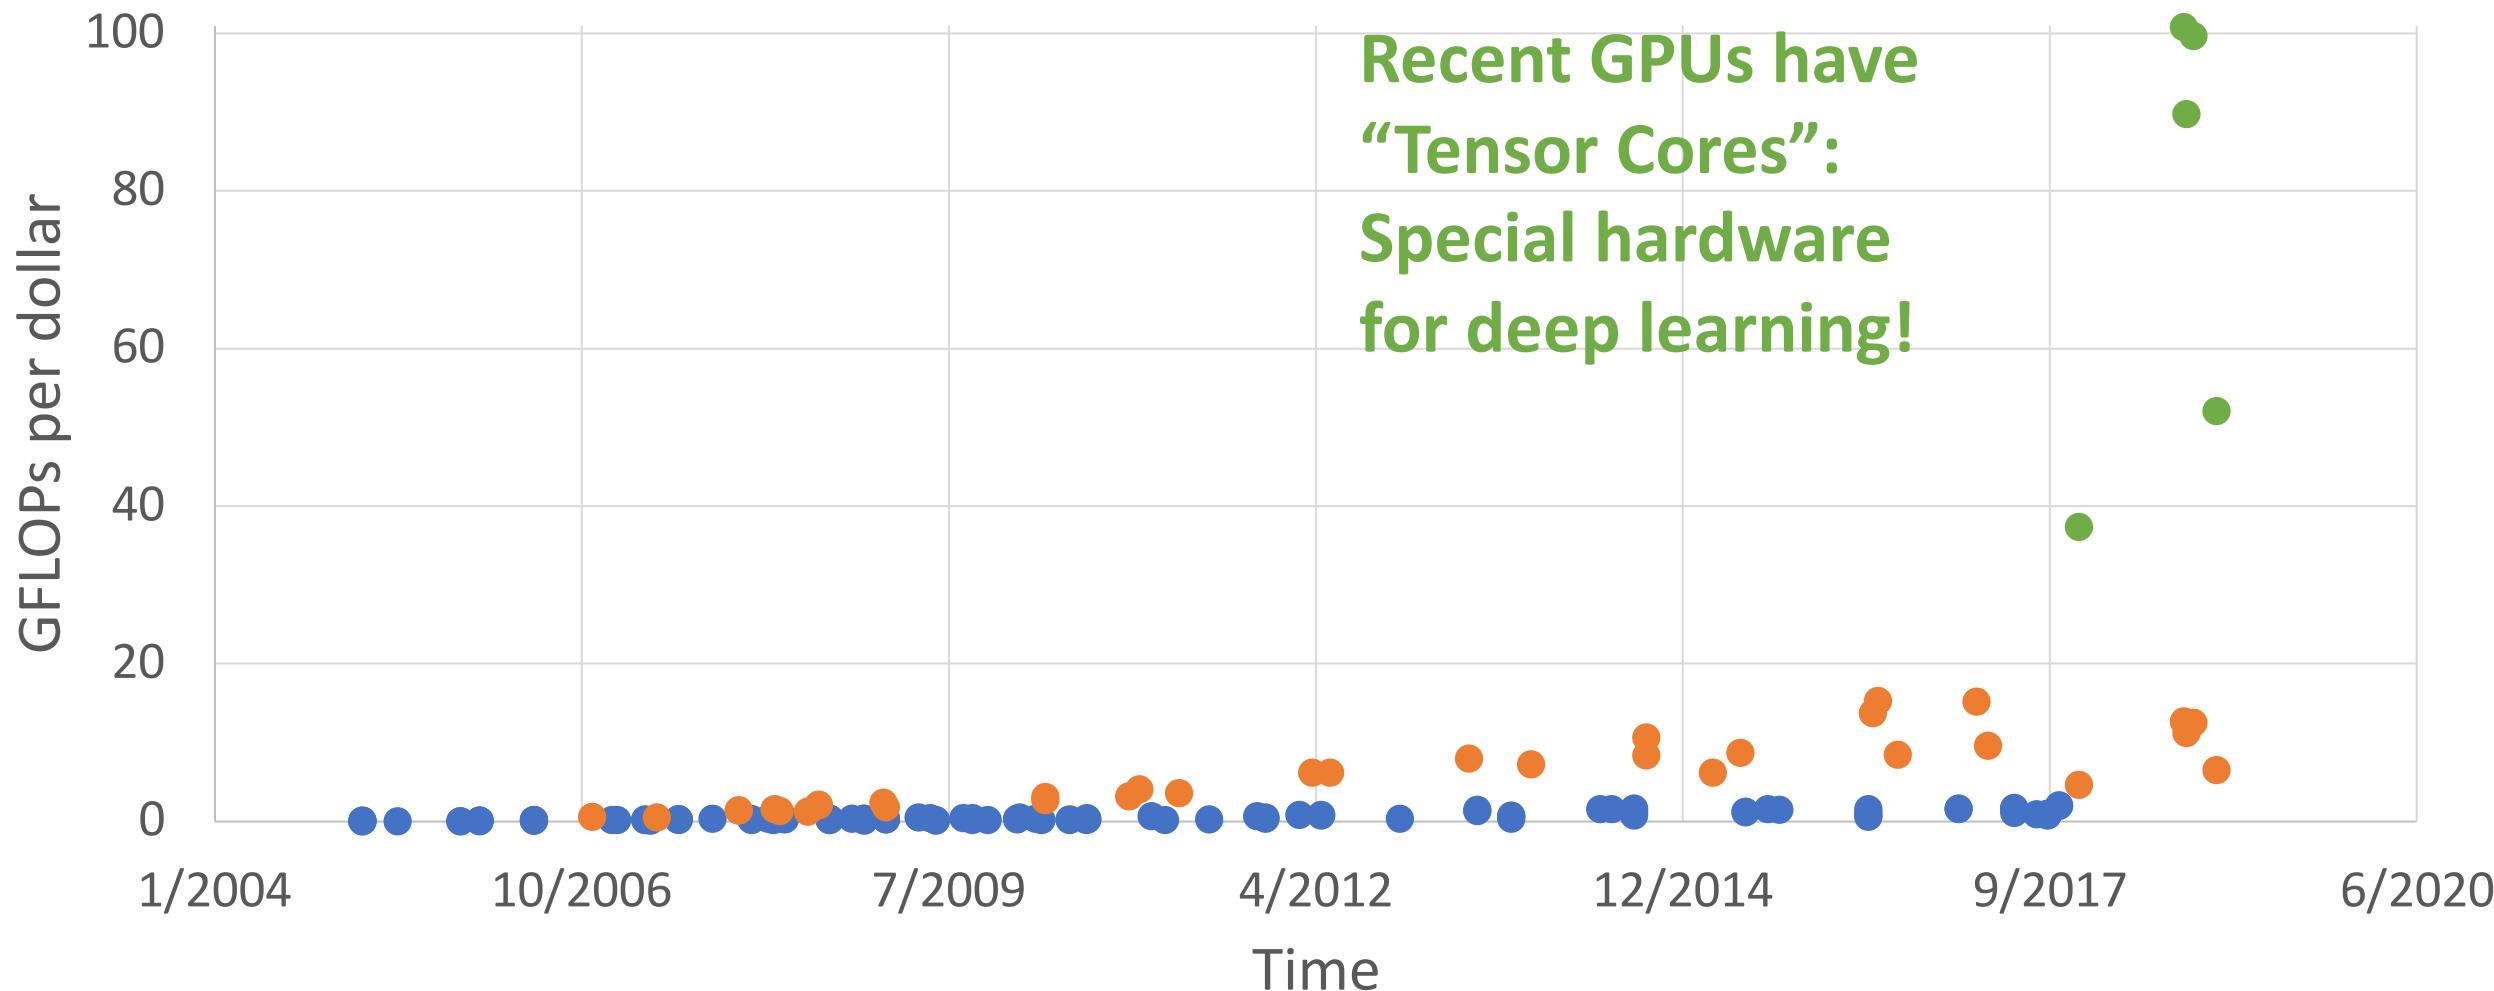
GFLOPs per Dollar

● CPU ● GPU FP32



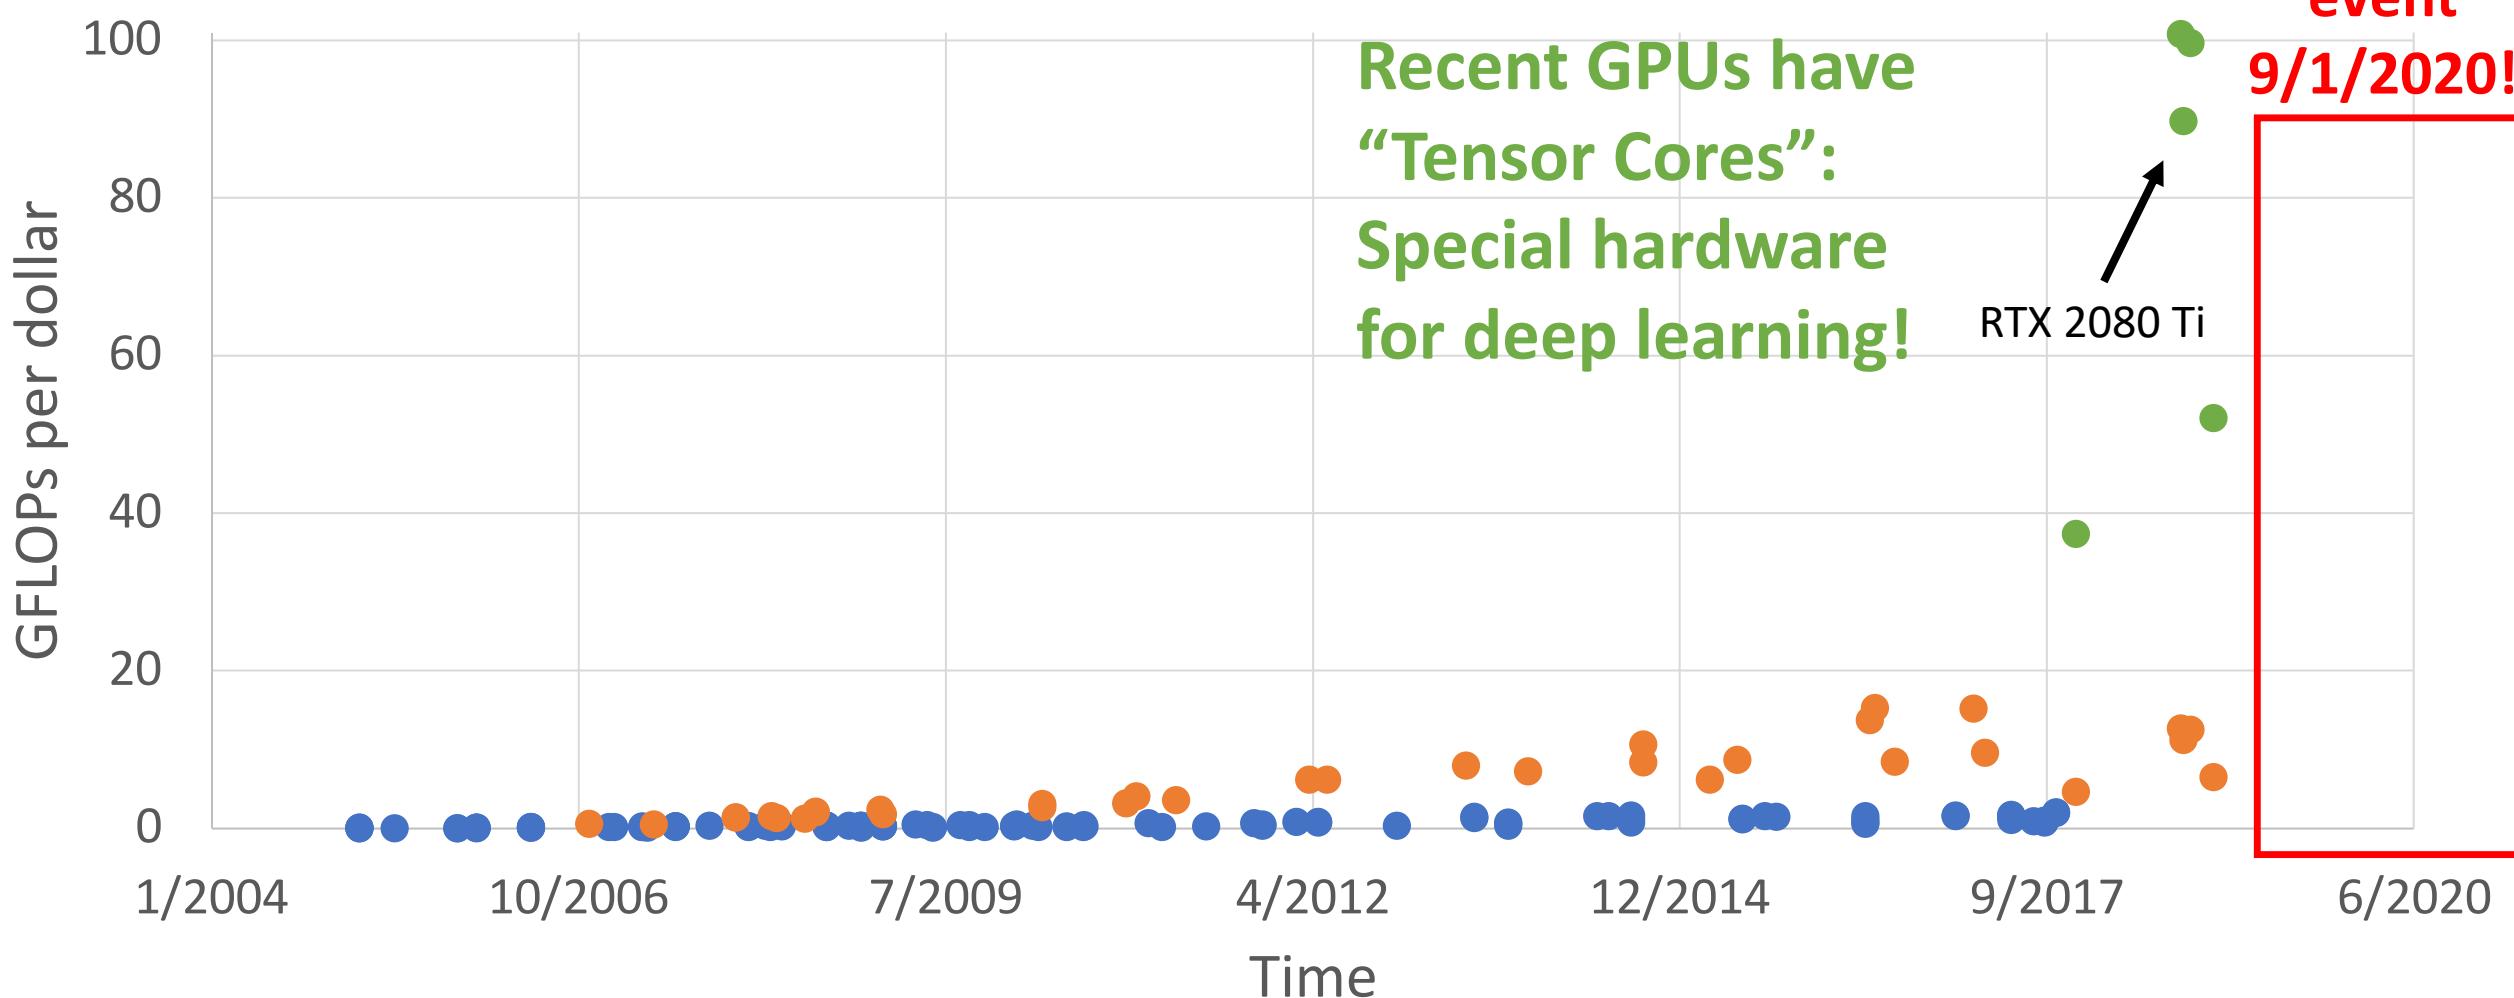
GFLOPs per Dollar

● CPU ● GPU FP32 ● GPU Tensor Core

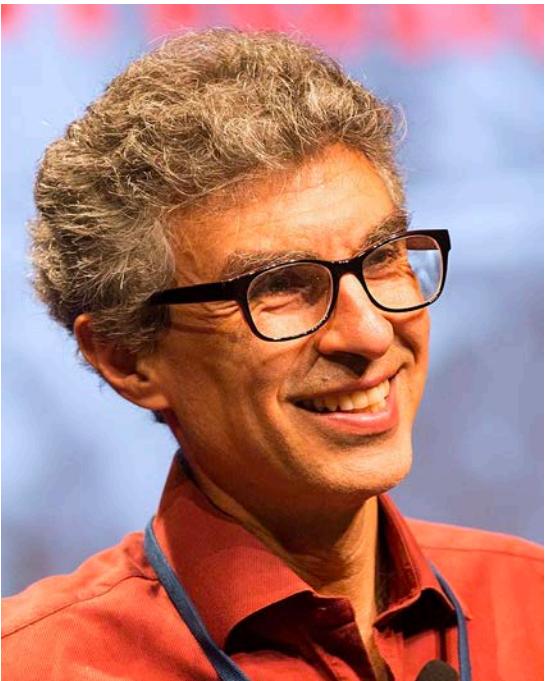


GFLOPs per Dollar

● CPU ● GPU FP32 ● GPU Tensor Core



2018 Turing Award



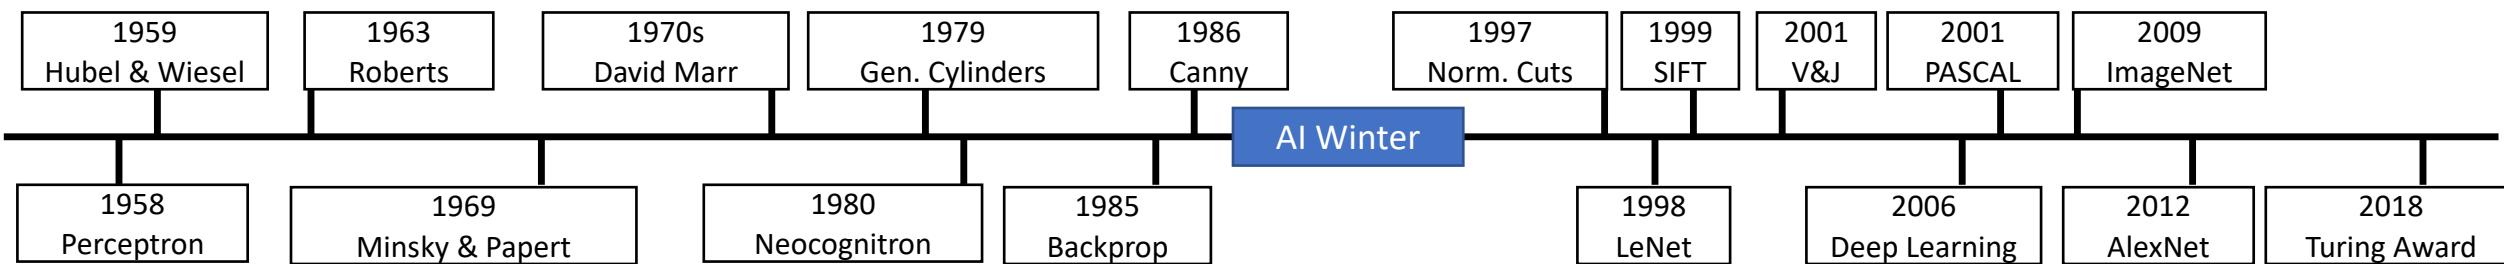
Yoshua Bengio



Geoffrey Hinton



Yann LeCun



Despite our success, computer vision still has a long way to go...

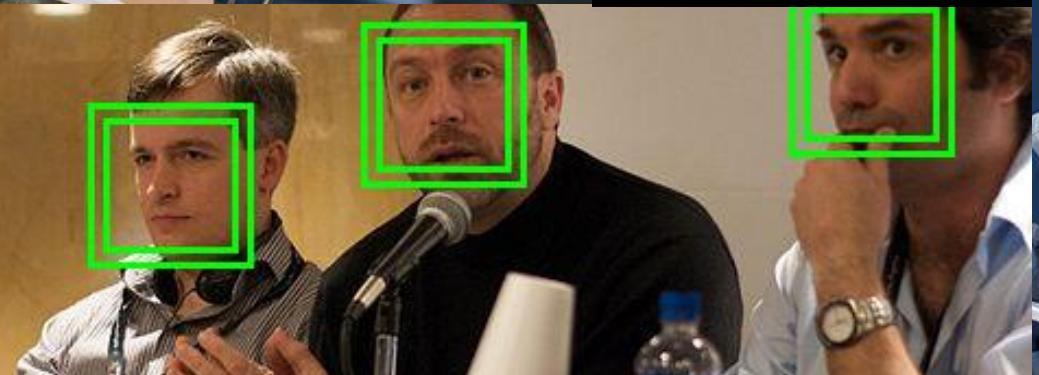
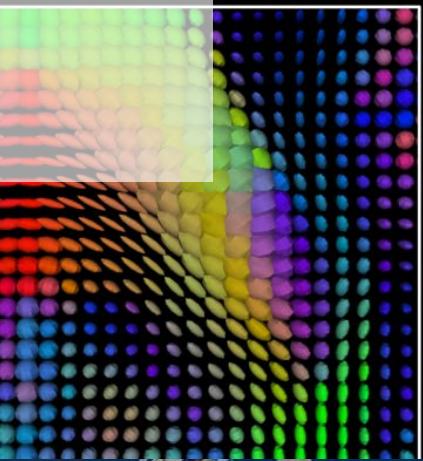


This image is copyright-free United States government work

Example credit:
Andrej Karpathy

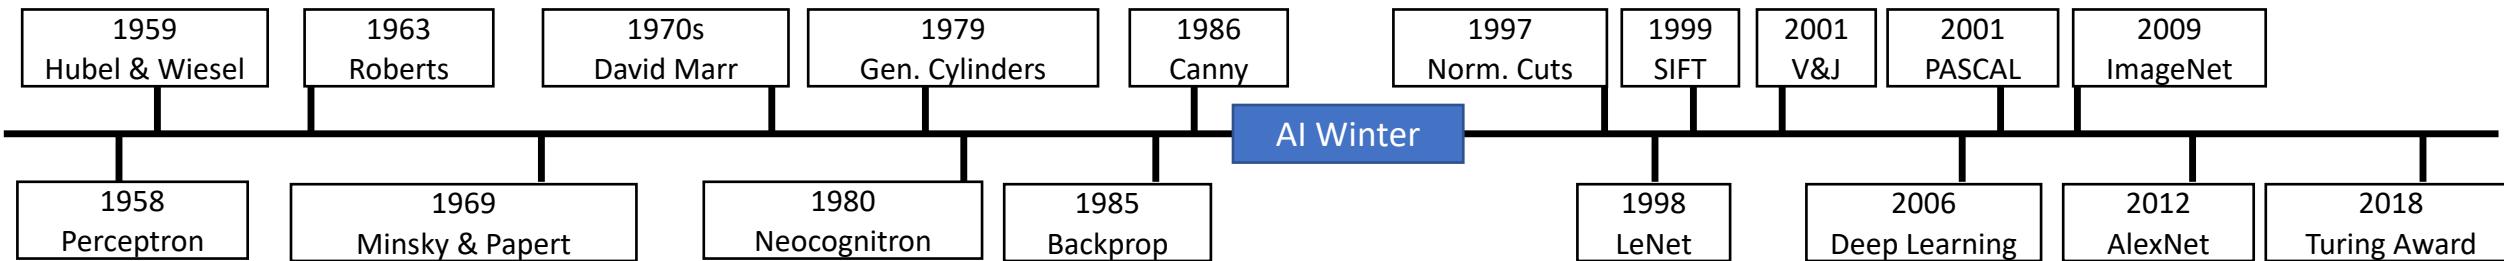
Computer Vision Technology

Can Better Our Lives



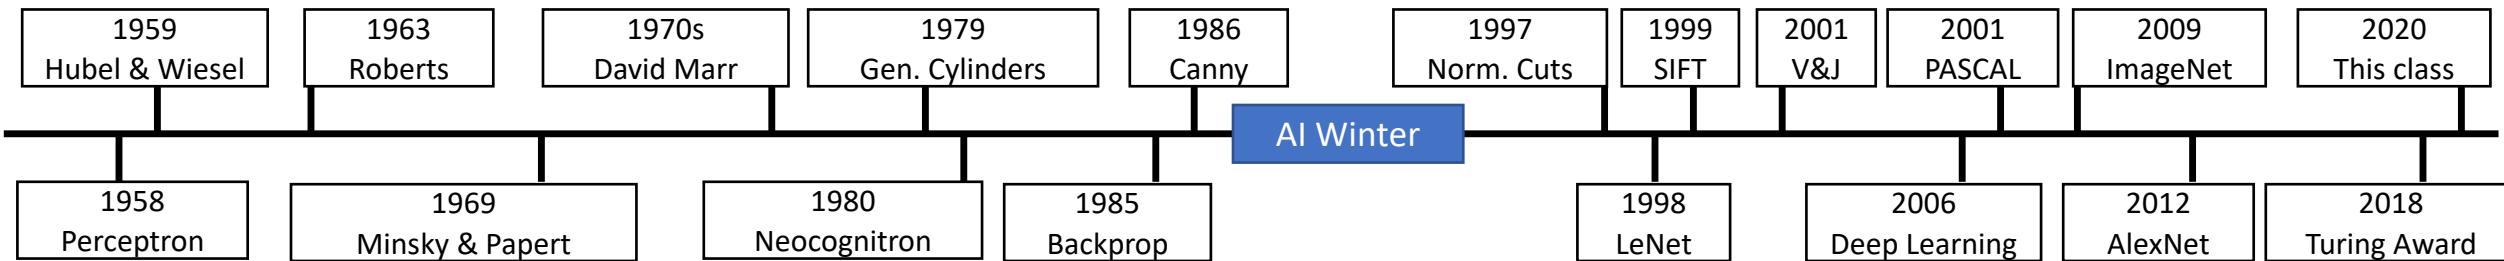
Today's Agenda

- A brief history of computer vision and deep learning
- Course overview and logistics



Today's Agenda

- A brief history of computer vision and deep learning
- Course overview and logistics



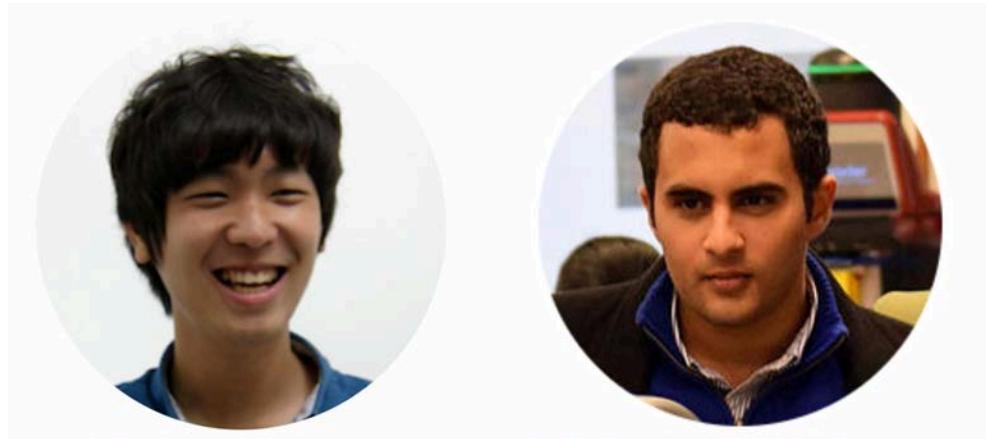
Course Staff

Instructor



Justin Johnson
Assistant Professor, CSE

Graduate Student Instructors



Yunseok Jang



Mohamed El Banani



Danish Syed



Yashmeet Gambhir

How to contact us

- Course Website: <https://web.eecs.umich.edu/~justincj/teaching/eecs498/>
 - Syllabus, schedule, assignments, slides, lecture videos, etc
- Piazza: <https://piazza.com/class/ke3a8m6u5wx647>
 - (Almost) all questions about the course should go here!
 - We will also use Piazza to communicate with you
 - Use private questions if you want to post code
- EECS Autograder:
 - For turning in homework assignments
 - Still working out details, will update soon
- [Google Calendar](#): For office hours (starting next week)
- Email: Only for sensitive, confidential issues

Course Website: Check the Schedule!



EECS 498-007 / 598-005
Deep Learning for Computer Vision
Fall 2020

Schedule

Lectures will be Mondays and Wednesdays 1:30 - 3pm on Zoom. Attendance is not required. Recordings will be posted after each lecture in case you are unable to attend the scheduled time.

Some lectures have reading drawn from the course notes of [Stanford CS 231n](#), written by [Andrej Karpathy](#).

Some lectures have optional reading from the book *Deep Learning* by Ian Goodfellow, Yoshua Bengio, and Aaron Courville (GBC for short). The entire text of the book is [available for free online](#) so you don't need to buy a copy.

Event	Date	Description	Course Materials
Lecture 1	Monday August 31	Course Introduction Computer vision overview Historical context Course logistics	[FA2019 slides] [FA2019 video] [Python tutorial] [GBC Sec 1.2] [GBC Sec 6.6]
Lecture 2	Wednesday September 2	Image Classification Data-driven approach K-Nearest Neighbor Hyperparameters Cross-validation	[FA2019 slides] [FA2019 video] [231n Image Classification]
	Monday September 7	No class Labor Day	

<https://web.eecs.umich.edu/~justincj/teaching/eecs498/FA2020/schedule.html>

Piazza Etiquette

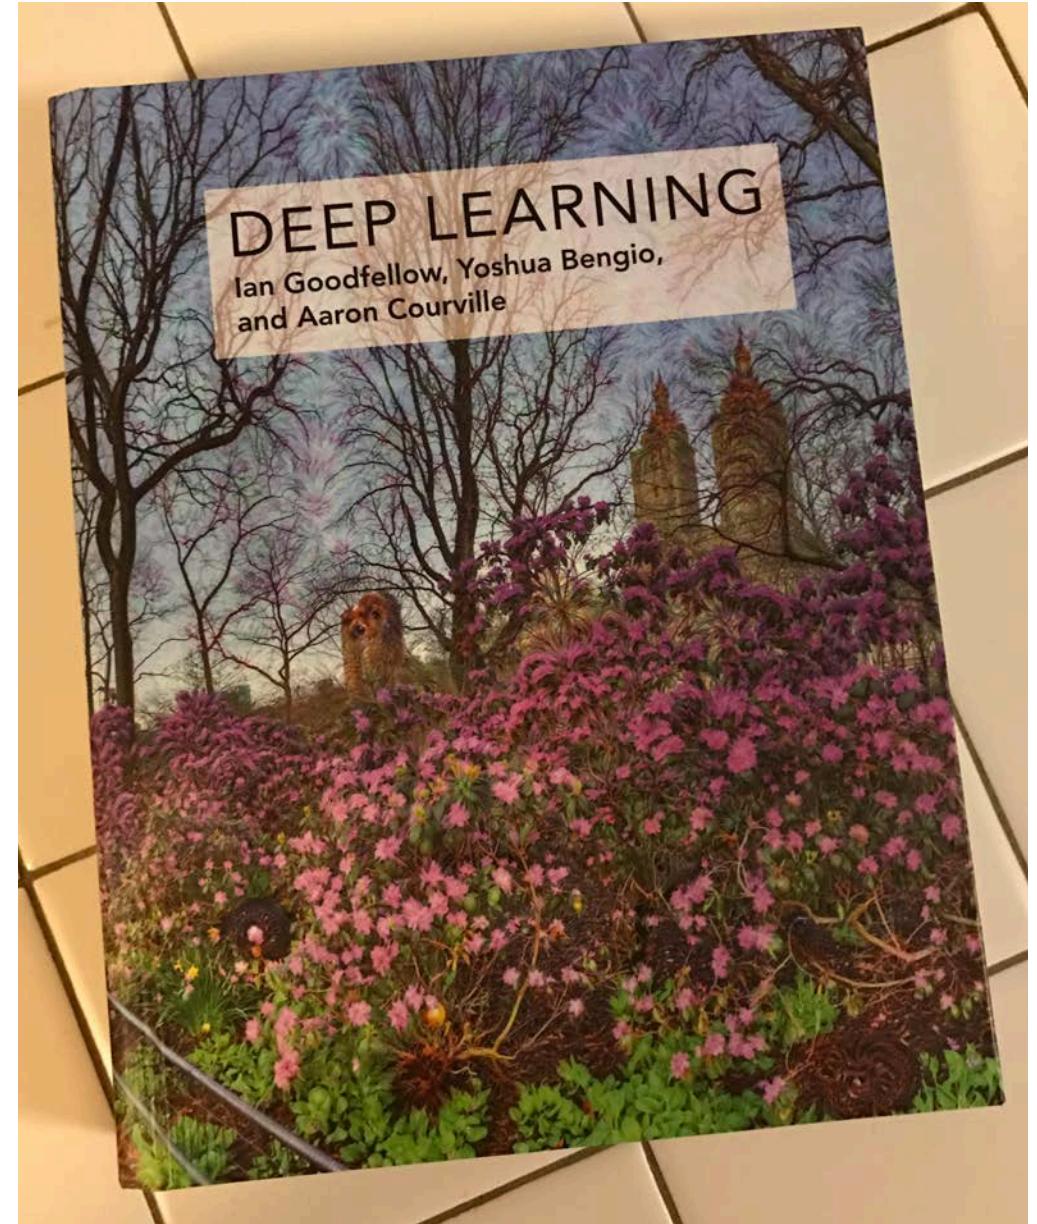
- Post only short snippets of code (< 20-30 lines)
- Ask a specific, concrete question
- Explain what you have tried so far, and what happened
- See StackOverflow guide on asking good questions:
<https://stackoverflow.com/help/how-to-ask>

Piazza Etiquette

- Post only short snippets of code (< 20-30 lines)
- Ask a specific, concrete question
- Explain what you have tried so far, and what happened
- See StackOverflow guide on asking good questions:
<https://stackoverflow.com/help/how-to-ask>
- Don't expect an answer within 30 minutes of posting
- Monday – Friday, 10am – 6pm EST we'll try to answer within 2 hours
- Other times, we'll try to answer within 12 hours

Optional Textbook

- [Deep Learning](#) by Goodfellow, Bengio, and Courville
- [Free online](#)



Course Content and Grading

- 6 programming assignments (A1 10%, A2-A6 12%)
 - Homework assignments will use Python, PyTorch, and Google Colab
- Midterm Exam (30%)
- Late policy
 - 3 free late days to use on assignments
 - Once free late days are exhausted, 25% penalty per day

Collaboration Policy

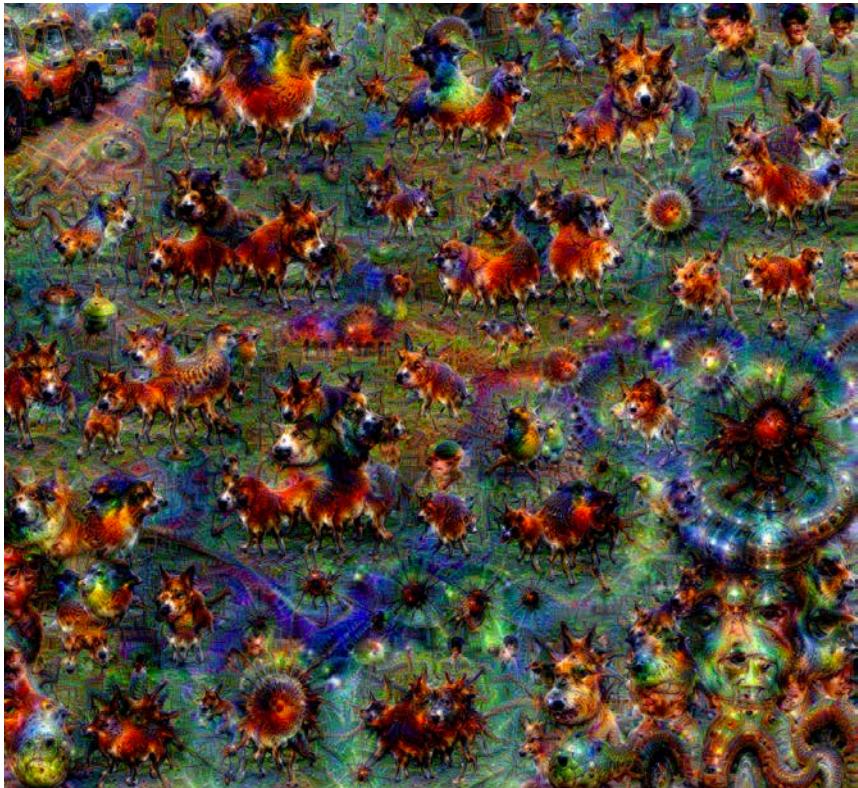
- **Rule 1:** Don't look at solutions or code that are not your own; everything you submit should be your own work
- **Rule 2:** Don't share your solution code with others; however discussing ideas or general strategies is fine and encouraged
- **Rule 3:** Indicate in your submissions anyone you worked with
- Turning in something late / incomplete is better than violating the honor code

Course Philosophy

- Thorough and Detailed.
 - This not “Learn PyTorch in 90 days”, nor “Deep Learning in 10 lines of code”
 - Understand how to write from scratch, debug, and train convolutional and other types of deep neural networks
 - We prefer to write from scratch, rather than rely on existing implementations
- Practical
 - Focus on practical techniques for training and debugging neural networks
 - Will use state-of-the-art software tools like PyTorch and TensorFlow
- State of the art
 - Most material we cover is research published in the last 5 years

Course Philosophy

- Will also cover some fun topics:
 - Image captioning
 - DeepDream, Artistic Style Transfer



Course Structure

- First half: Fundamentals
 - Details of how to implement and train different types of networks
 - Fully-connected networks, convolutional networks, recurrent networks
 - How to train and debug, very detailed
- Second half: Applications and “Researchy” topics
 - Object detection, image segmentation, 3D vision, videos
 - Attention, Transformers
 - Vision and Language
 - Generative models: GANs, VAEs, etc
 - Guest Lectures from subject-matter experts
 - Less detailed: provide overview and references, but skip some details

First homework assignment

- Will be released by tomorrow
- Due Friday 9/11/2020
- Next lecture will be enough to complete it

Next time: Image Classification